

# **Learning-based Cardinality/Cost Estimation**

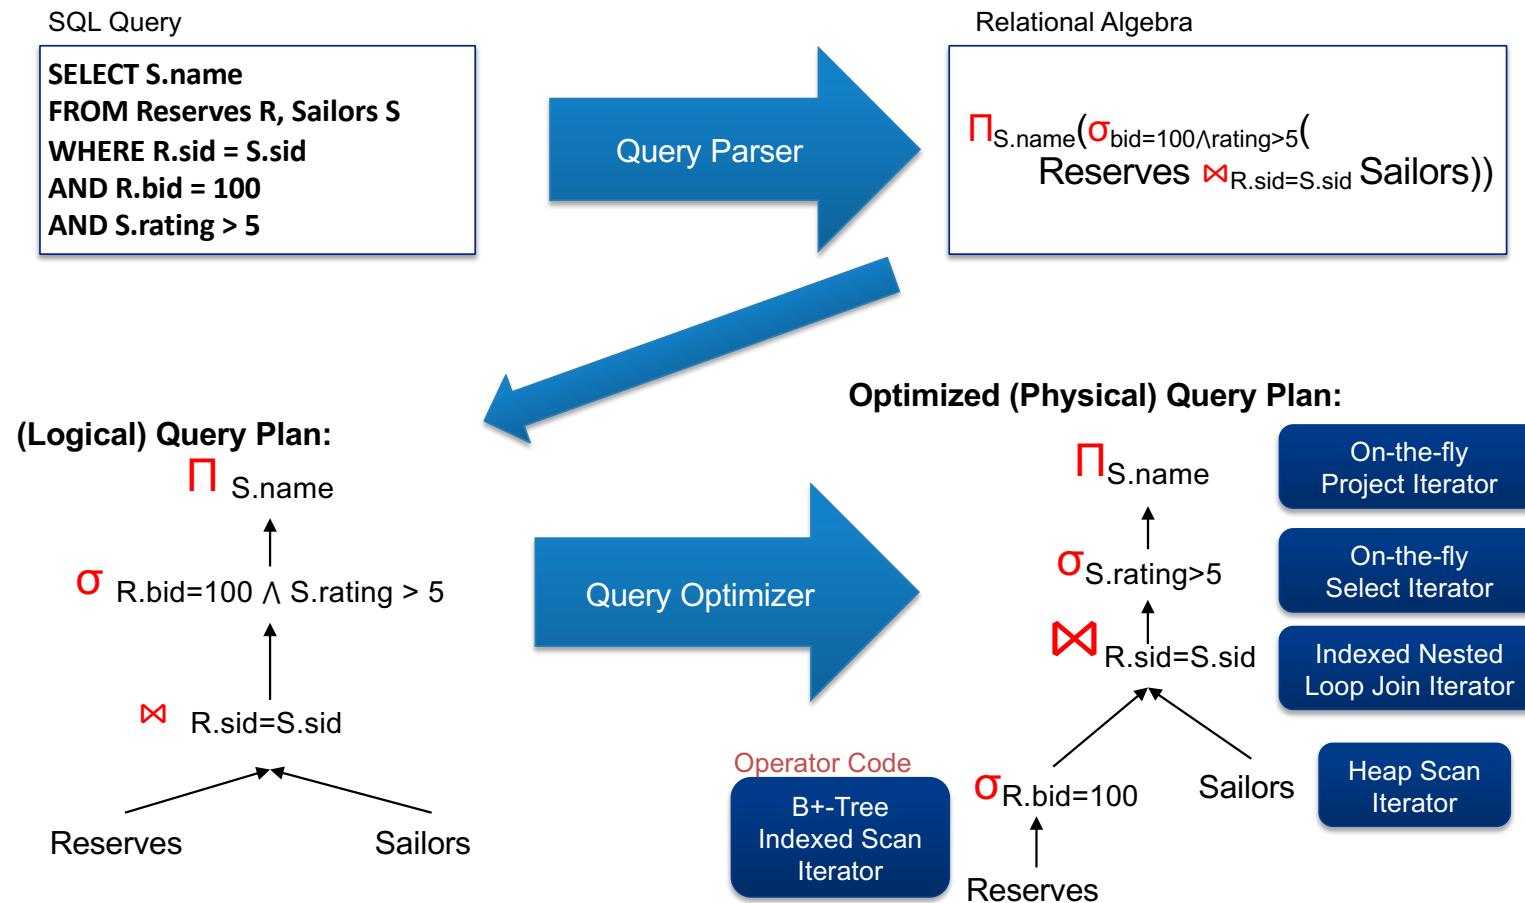
**Guoliang Li**

**Tsinghua University**



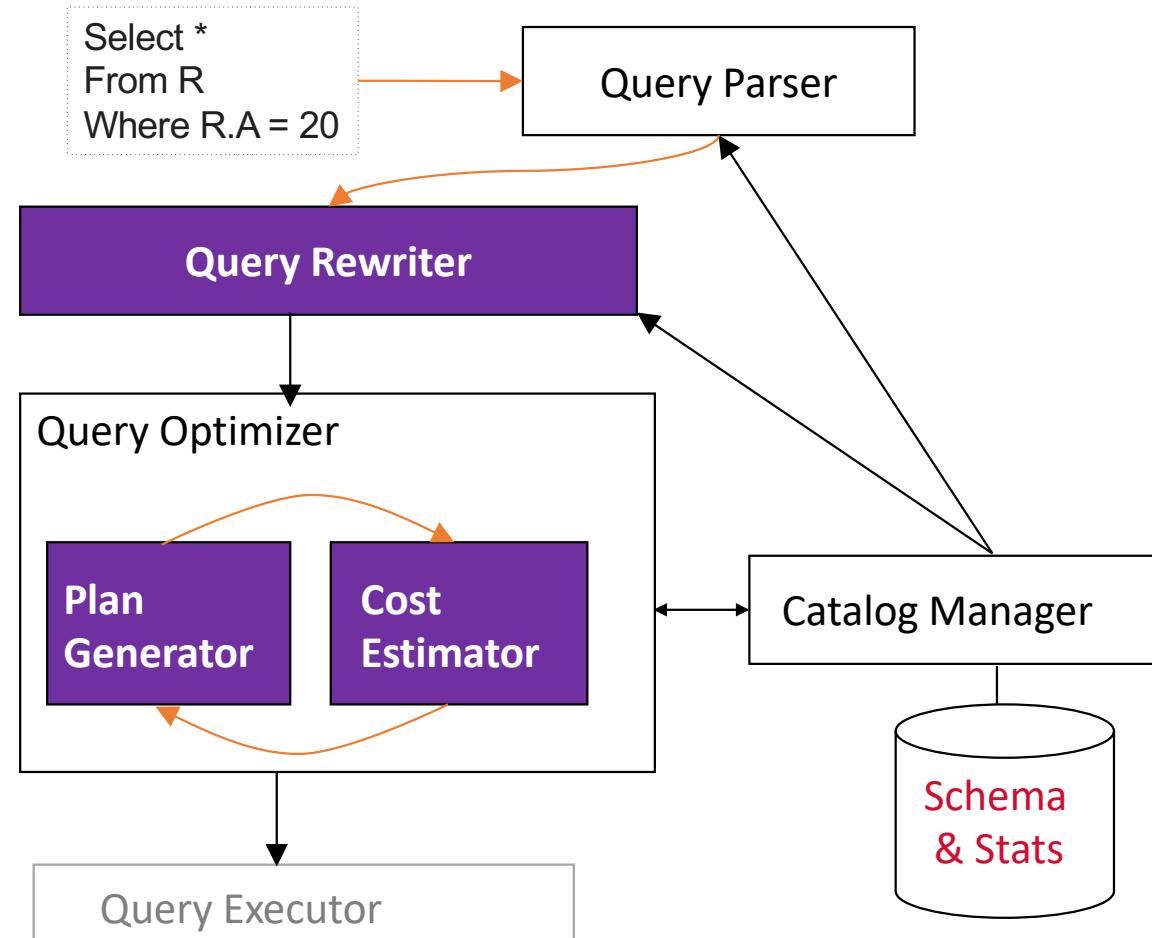


# Query Optimizer



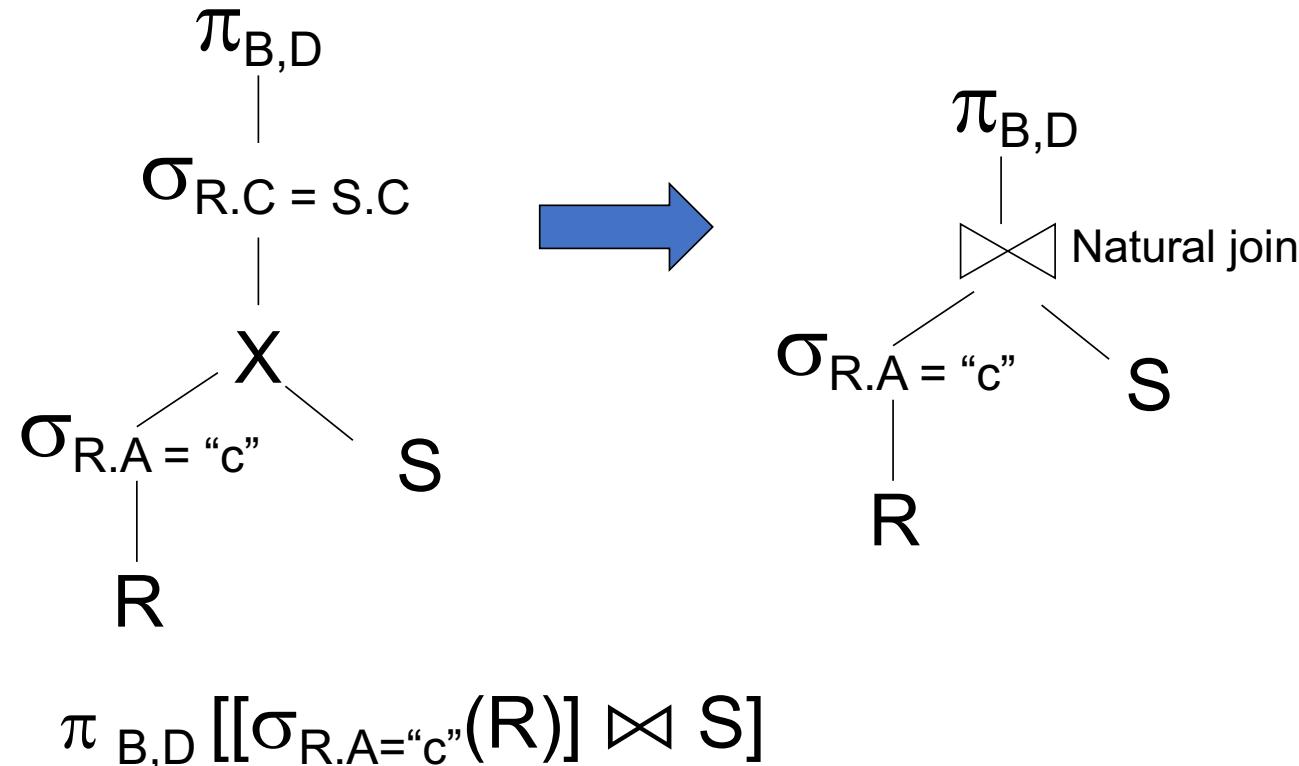


# Query Optimizer



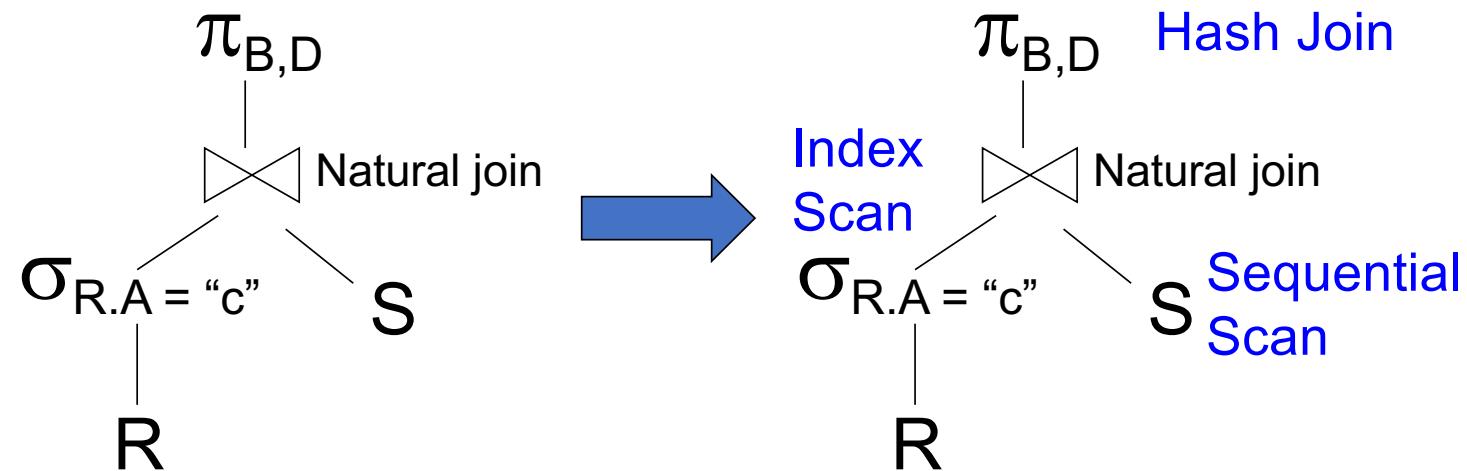


# Logical Optimization – Query Rewrite





# Physical Optimization



$$\pi_{B,D} [\sigma_{R.A = "c"}(R) \bowtie S]$$



# Cardinality Estimation: Selection



Selectivity Factor (SF) = Cardinality / #tuples

Assumptions:

- Uniformity
  - $\sigma_{A=c}(R) \rightarrow SF = 1/V(R,A)$
  - $\sigma_{A < c}(R) \rightarrow SF = (c - Low(R, A)) / (High(R, A) - Low(R, A))$
- Independence
  - Cond1 and Cond2  $\rightarrow SF = SF(Cond1) * SF(Cond2)$
  - Cond1 or Cond2  $\rightarrow SF = SF(Cond1) + SF(Cond2) - SF(Cond1) * SF(Cond2)$
  - Not Cond1  $\rightarrow SF = 1 - SF(Cond1)$
- Containment of values
  - $R_{\bowtie A=B} S \rightarrow SF = 1 / \max(V(R,A), V(S,B))$
- Preservation of values
  - $V(R_{\bowtie A=B} S, C) = V(R, C)$



# Cardinality Estimation: Selection

```
Q = SELECT list  
      FROM R1, ..., Rn  
      WHERE cond1 AND cond2 AND ... AND condk
```

- Estimate the number of results of Q:  $T(Q)$
- Obtain number of tuples in each table:  $T(R1), T(R2), \dots, T(Rn)$
- Also need the **selectivity** of each condition
  - Selectivity factor (SF) of selection and joins
  - $SF(R1.A=v)=T(R1)/V(R1,A)$
  - $SF(R1.A=R2.B)=T(R1) T(R2)/ \max(V(R1,A), V(R2,B))$
  - e.g., selectivity(A=3) = 0.01
  - e.g., selective (R1.A=R2.B) = 0.001

$T(Q) = T(R1) \times \dots \times T(Rn) \times SF(cond1) \times \dots \times SF(condk)$   
Remark:  $T(Q) \leq T(R1) \times T(R2) \times \dots \times T(Rn)$



# Cardinality Estimation: Predicate

- The **selectivity (sel)** of a predicate **P** is the fraction/probability of tuples that qualify.
- Formula depends on type of predicate:
  - Equality(=):  $\text{sel}(P(c=x)) = \text{count}(c=x) / \text{count}(\text{all})$
  - Range( $\geq$ ):  $\text{sel}(P(c \geq x)) = (\text{max} - x + 1) / (\text{max} - \text{min} + 1)$
  - Negation ( $\neq$ ) :  $\text{sel}(P(c \neq x)) = 1 - \text{sel}(P(c=x))$
  - Conjunction (and)
    - Independent assumption
      - $\text{sel}(P1 \ \& \ P2) = \text{sel}(P1) * \text{sel}(P2)$
  - Disjunction (or)
    - $\text{sel}(P1 \ \text{or} \ P2) = \text{sel}(P1) + \text{sel}(P2) - \text{sel}(P1) * \text{sel}(P2)$



# Cardinality Estimation: Join

$R \bowtie_{A=B} S$ : Selectivity =  $1 / \max(V(R,A), V(S,A))$

Q= **SELECT \* FROM R, S, T**  
**WHERE R.B=S.B and S.C=T.C and R.A=40**

$T(R) = 30k$ ,  $T(S) = 200k$ ,  $T(T) = 10k$

Selectivity of  $R.B = S.B$  is  $1/3$

Selectivity of  $S.C = T.C$  is  $1/10$

Selectivity of  $R.A = 40$  is  $1/200$

$R(A,B)$

$S(B,C)$

$T(C,D)$

$$T(Q) = 30k * 200k * 10k * 1/3 * 1/10 * 1/200 = 10^{10}$$

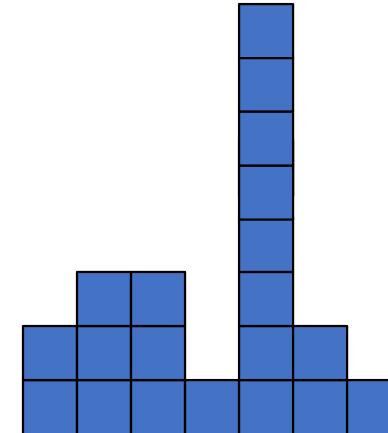


# Histograms

- For better estimation, use a histogram

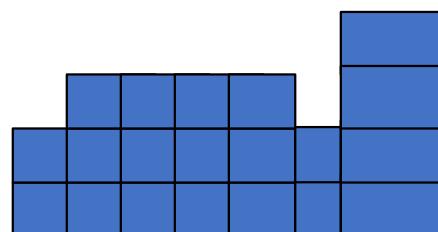
*equiwidth*

No. of Values	2	3	3	1	8	2	1
Value	0-.99	1-1.99	2-2.99	3-3.99	4-4.99	5-5.99	6-6.99



*equidepth*

No. of Values	2	3	3	3	3	2	4
Value	0-.99	1-1.99	2-2.99	3-4.05	4.06-4.67	4.68-4.99	5-6.99



Note: 10-bucket equidepth histogram divides the data into *deciles*  
- akin to quantiles, median, etc.  
Common trick: “end-biased” histogram  
- very frequent values in their own buckets



# Sketch

- How to count the number of values in a column?
  - E.g., Age = 20?
- Sketch: a cost-model can replace histograms with sketches to improve its selectivity estimate accuracy.
- Probabilistic data structures that generate approximate statistics about a data set.
- Most common examples:
  - Count-Min Sketch (1988): Approximate frequency count of elements in a set.
  - HyperLogLog (2007): Approximate the number of distinct elements in a set.



# Sketch

- Given a column with a set of values, build a hash function  $H$ , and a sketch  $M$  with  $w$  entries.
  - $M[i]$  is initialized as 0 for  $0 \leq i \leq w-1$
  - For each value  $v$  in the set,
    - $M[H[v] \% w] ++;$
- Given a value  $x$ , use  $M[H[x] \% w]$  to estimate its account.

Overestimate! Because of collisions!



# Count-min Sketch

- How about use  $d$  hash functions to reduce collisions?
- A matrix  $M$  with  $d$  rows and  $w$  columns, initialized with 0 for each cell value;  $d$  hash functions
- For each value  $v$ , each hash function  $h_i$ :
  - $M[i][h_i(v)] = M[i][h_i(v)] + 1$ ;  $h_i(v)$  in  $[0, w]$
- Given  $x$ , the frequency  $f(x)$  can be estimated as
  - $f(x) = \min_{i \text{ in } [0, d-1]} M[i][h_i(v)]$

	0	1	2	3	4	...	w-1
$h_0(v)$	1	0	0	2	0	0	1
$h_1(v)$	1	0	0	1	0	0	0
...	0	0	0	0	2	0	0
$H_{d-1}(v)$	3	0	0	1	0	0	1



# Count-min Sketch

- $d=4$  hash functions,  $w=7$  columns
- Given values  $\{2, 3, 2, 4, 3, 2, 5\}$ 
  - $h_0(v) = v \% w; h_1(v) = v^2 \% w; h_2(v) = (2v+1) \% w; h_3(v) = (3v^2+1) \% w;$
- Add 2,

	0	1	2	3	4	5	6
$h_0(v)$	0	0	1	0	0	0	0
$h_1(v)$	0	0	0	0	1	0	0
$h_2(v)$	0	0	0	0	0	1	0
$h_3(v)$	0	0	0	0	0	0	1



# Count-min Sketch

- $d=4$  hash functions,  $w=7$  columns
- Given values  $\{2, 3, 2, 4, 3, 2, 5\}$ 
  - $h_0(v) = v \% w; h_1(v) = v^2 \% w; h_2(v) = (2v+1) \% w; h_3(v) = (3v^2+1) \% w;$
- Add 2, 3

	0	1	2	3	4	5	6
$h_0(v)$	0	0	1	1	0	0	0
$h_1(v)$	0	0	1	0	1	0	0
$h_2(v)$	1	0	0	0	0	1	0
$h_3(v)$	1	0	0	0	0	0	1



# Count-min Sketch

- $d=4$  hash functions,  $w=7$  columns
- Given values  $\{2, 3, 2, 4, 3, 2, 5\}$ 
  - $h_0(v) = v \% w; h_1(v) = v^2 \% w; h_2(v) = (2v+1) \% w; h_3(v) = (3v^2+1) \% w;$
- Add 2, 3, 2

	0	1	2	3	4	5	6
$h_0(v)$	0	0	2	1	0	0	0
$h_1(v)$	0	0	1	0	2	0	0
$h_2(v)$	1	0	0	0	0	2	0
$h_3(v)$	1	0	0	0	0	0	2



# Count-min Sketch

- $d=4$  hash functions,  $w=7$  columns
- Given values  $\{2, 3, 2, 4, 3, 2, 5\}$ 
  - $h_0(v) = v \% w; h_1(v) = v^2 \% w; h_2(v) = (2v+1) \% w; h_3(v) = (3v^2+1) \% w;$
- Add 2, 3, 2, 4

	0	1	2	3	4	5	6
$h_0(v)$	0	0	2	1	1	0	0
$h_1(v)$	0	0	1+1	0	2	0	0
$h_2(v)$	1	0	1	0	0	2	0
$h_3(v)$	1+1	0	0	0	0	0	2



# Count-min Sketch

- $d=4$  hash functions,  $w=7$  columns
- Given values  $\{2, 3, 2, 4, 5\}$ 
  - $h_0(v)=v \% w; h_1(v)=v^2 \% w; h_2(v)=(2v+1) \% w; h_3(v)=(3v^2+1) \% w;$
- Add  $2, 3, 2, 4, \underline{5}$

	0	1	2	3	4	5	6
$h_0(v)$	0	0	2	1	1	1	0
$h_1(v)$	0	0	1+1	0	2+1	0	0
$h_2(v)$	1	0	1	1	0	2	0
$h_3(v)$	1+1	0	0	0	0	0	2+1

$\text{count}(2)=2$ ;  $\text{count}(3)=1$ ;  $\text{count}(4)=1$ ;  $\text{count}(5)=1$



# Loglog

Crucial insight: suppose we have a perfect hash function  $h$  taking an integer from  $[1,r]$  and reporting an integer  $[0, n)$

$$h(x) \in [0, n) \text{ for } x \in [1, r]$$

The probability that the hash contains:

0 leading zeroes:  $1/2$

1 leading zero:  $1/4$

2 leading zeroes:  $1/8$

3 leading zeroes:  $1/16$

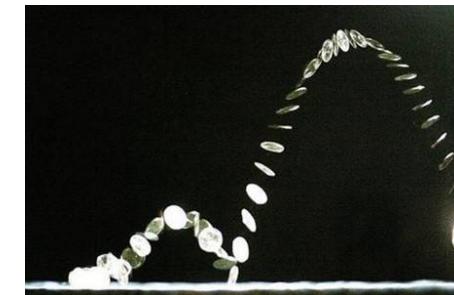
...

$w < \log n$  leading zeroes  $1/2^w$



# LogLog

- Estimate the number of distinct values in a column
- Linear Counting: inefficient
- Hash: large space
- Log Counting:  $n=2^{\max(\text{leading } 0\text{s})}$
- $dv=0$ : initial distinct number
- For each element  $v$  in the column
  - Hash( $v$ ) to 0/1 values
  - $C_0=\text{count leading } 0\text{s}$
  - $dv=\max(dv, 2^{C_0})$





# Loglog

- Crucial insight: maintaining the **largest number of leading zeroes** across all hashes allows us to get a (very) rough estimate of the number of distinct elements.
- take multiple independent hashes for each element, average them out?

$$\text{Estimate: } 2^3 = 8$$

Actual: 10

7 → 11101100

6 → 11010101

15 → 10110100

8 → 11100110

15 → 10110100

2 → 00100010

2 → 00100010

9 → 01100100

11 → 00011000

8 → 11100110

14 → 10110111

16 → 01001101

12 → 00110110

6 → 11010101

2 → 00100010



# Loglog

$$\sum 2^w$$

Estimate:  $2^4 + 2^2 + 2^0 + 2^1 = 23$   
Actual: 10

## 00 stream

2 → 00100010  
2 → 00100010  
11 → 00000010  
12 → 00110110  
2 → 00100010

4

## 10 stream

15 → 10110100  
15 → 10110100  
14 → 10110111

0

## 11 stream

7 → 11101100  
6 → 11010101  
8 → 11100110  
8 → 11100110  
6 → 11010101

2

1

## 01 stream

9 → 01100100  
16 → 01001101



# Loglog

$m \times 2^{(\sum w)/m}$

Estimate:  $4 \cdot 2^{(4+2+0+1)/4} \approx 13.45$   
Actual: 10

## 00 stream

2 → 00100010  
2 → 00100010  
11 → 00000010  
12 → 00110110  
2 → 00100010

4

## 10 stream

15 → 10110100  
15 → 10110100  
14 → 10110111

0

## 11 stream

7 → 11101100  
6 → 11010101  
8 → 11100110  
8 → 11100110  
6 → 11010101

2

1

## 01 stream

9 → 01100100  
16 → 01001101



# HyperLoglog

LogLog uses the arithmetic mean

$$\alpha_m \times m \times 2^{(\sum w)/m}$$

HyperLogLog uses an alternative, the harmonic mean

$$\alpha_m \times m \times m / (\sum 2^{-w})$$

where  $\alpha_m$  is a constant.

$$\alpha_m = \begin{cases} 0.673, & m = 16 \\ 0.697, & m = 32 \\ 0.709, & m = 64 \\ \frac{0.7213}{1 + \frac{1.079}{m}}, & m \geq 128 \end{cases}$$



# Loglog

Thus the LogLog algorithm can be succinctly described. We need a hash function  $h$  from  $[1,r]$  to  $[0, n)$

1. Initialize an array  $w$  with size  $m=2^b$ . Let  $w_i = 0$  for all  $i$
2. For each element  $x$ , compute its hash  $h(x)$ 
  - Split  $h(x)$  to its first  $b$  bits and remaining  $m-b$  bits.
  - Let  $c$  be the number represented by the first  $b$  bits and  $w_0$  be the number of leading zeroes in the  $m-b$  bits.
  - Set  $w_c = \max(w_c, w_0)$
3. Output

$$\alpha_m \times m \times 2^{(\sum w_c)/m}$$

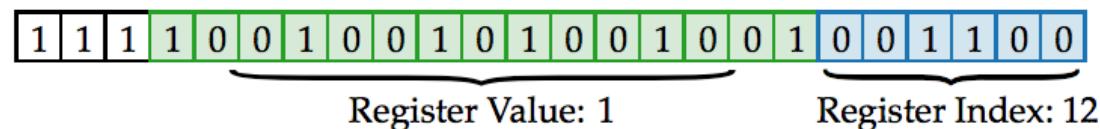


# HyperLogLog

- Estimate the number of distinct values in a column
- $m$  buckets,  $\alpha_m$  : a constant;  $M_j$ : # of leading 0s
- LogLog:  $\alpha_m \cdot m \cdot 2^{\sum_{j=1}^m M_j/m}$
- HyperLoglog:  $\alpha_m \cdot m^2 \cdot \left( \sum_{j=1}^m 2^{-M[j]} \right)^{-1}$

Value: 10,492,800

Hash Value: 1,475,498,572



Register Values:

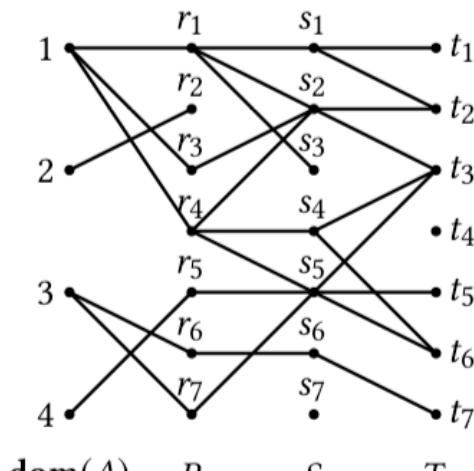
$m = 64$

0	0	0	0	0	0	0	1	0	0	0	1	0	1	0	0	0	0
2	0	0	0	0	0	0	0	0	2	0	0	3	0	2	0	0	0
0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	3	0	0	0	0	0	0	0	0	0	0	0	0	0	0

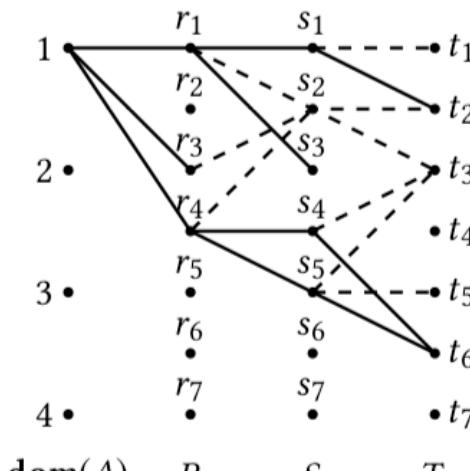


# Sampling

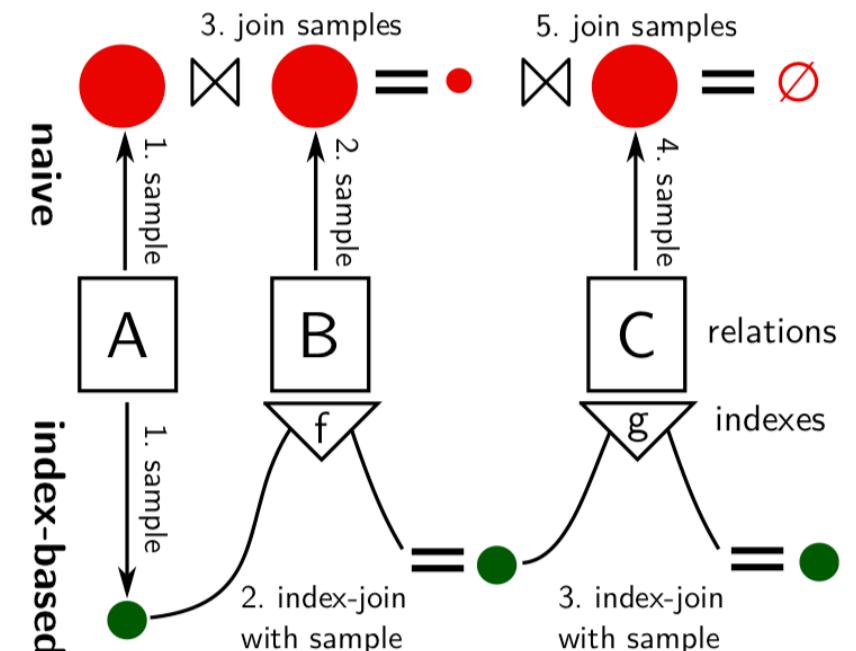
- Modern DBMSs also collect samples from tables to estimate selectivity.
- Update samples when the underlying tables changes significantly.



(a) Adding conceptual column.



(b) Distinct sampling from join.





# Traditional Cost Model Summary

- Sampling-based
- Histogram
- Sketch

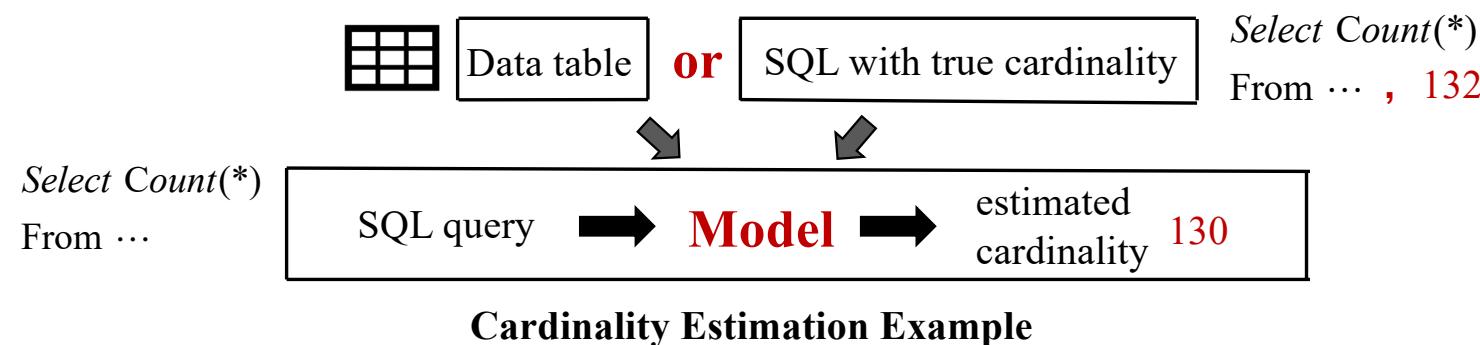
Algorithm	Ref.	Year	Observables [14]	Intuition [27]	Core method (Sec. 3)
FM	[15]	1985	Bit-pattern	Logarithmic hashing	Count trailing 1s
PCSA	[15]	1985	Bit-pattern	Logarithmic hashing	Count trailing 1s
AMS	[3]	1996	Bit-pattern	Logarithmic hashing	Count leading 0s
BJKST	[4]	2002	Order statistics	Bucket-based	Count leading 0s
LogLog	[11]	2003	Bit-pattern	Logarithmic hashing	Count leading 0s
SuperLogLog	[11]	2003	Bit-pattern	Logarithmic hashing	Count leading 0s
HyperLogLog	[14]	2008	Bit-pattern (order statistics)	Logarithmic hashing	Count leading 0s
HyperLogLog++	[24]	2013	Bit-pattern	Logarithmic hashing	Count leading 0s
MinCount	[21]	2005	Order statistics	Interval-based	k-th minimum value
AKMV	[7]	2007	Order statistics	Interval-based	k-th minimum value
LC	[32]	1990	No observable	Bucket-based	Linear synopses
BF	[28]	2010	No observable	Bucket-based	Linear synopses



# Cardinality Estimation

## □ Problem Formulation:

- **Cardinality:** The result size of a query.
- **Input:** A SQL Query.
- **Output:** An Estimated Cardinality.





# Traditional Cardinality Estimation

## □ Sampling

- **Core idea:** Estimating selectivity of target query by sampling.
- **Limitation:** Inference is slow and inaccurate when the amount of data is large.

## □ Histogram

- **Core idea:** Store the value distribution of each attribute, and calculate the selectivity according to the independence assumption.
- **Limitation:** Strong independence assumption makes it inaccurate when the data distribution is complex.
- **Sketch:** Estimate the number of distinct elements of a set.



# Regression Problems

**Database estimation problems** can be modeled as regression problems, which fit the high-dimension input variables into target features (e.g., cost, utility) and estimate the value of another variable.

- **Cardinality/Cost Estimation** aims to estimate the cardinality of a query and a regression model (e.g., deep learning model) can be used.
- **Index/View Benefit Estimation** aims to estimate the benefit of creating an index (or a view), and a regression model can be used to estimate the benefit.
- **Query Latency Estimation** aims to estimate the execution time of a query and a regression model can be used to estimate the performance based on query and concurrency features.



# Learned Cardinality Estimation

## □ Motivation

- Due to the attribute **value independence assumption** as well as the **assumption of uniform distribution**, traditional cardinality estimation methods tend to fail when the data distribution is complex.

## □ Challenge

- **Correlation** between data columns and columns.
- **Multi-table join** increases data volume and query types.

## □ Optimization Goal

- Accuracy, Inference Latency, Model Size, Training Cost



# Automatic Cardinality/Cost Estimation

## □ Motivation:

### □ One of the most challenging problems in databases

- Achilles Heel of modern query optimizers

### □ Traditional methods for cardinality estimation

- Sampling (on base tables or joins)
- Kernel-based Methods (Gaussian Model on Samples)
- Histogram (on single column or multiple columns)

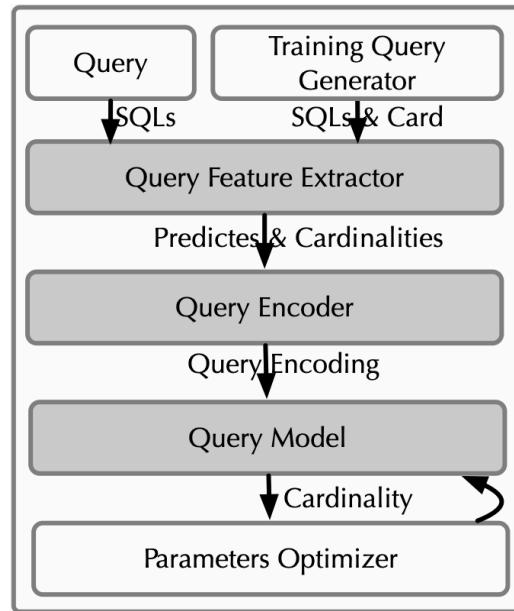
### □ Traditional cost models

- Data sketching/data histogram based methods
- Sampling based methods

Viktor Leis, Andrey Gubichev, Atanas Mirchev, Peter Boncz, Alfons Kemper, and Thomas Neumann. How good are query optimizers, really? In VLDB, 2015.



# Categories of Learned Cardinality Estimation

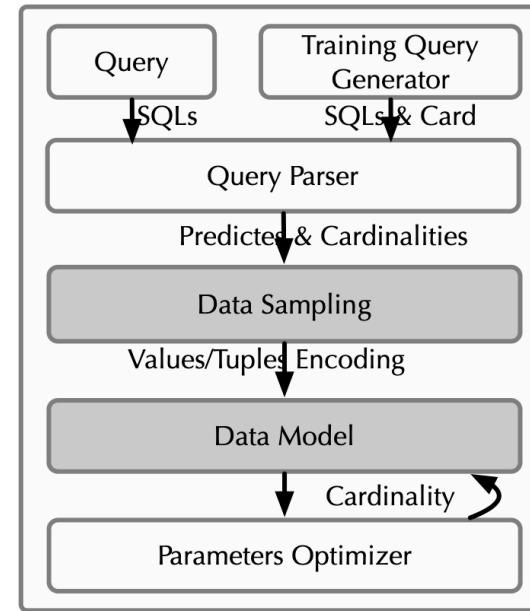


## (1) Supervised Query Methods

- Neural Network
- XGBoost
- Multi-set Convolutional network

## (2) Supervised Data Methods

- Kernel-density model
- Uniform mixture model
- Pre-training summarization model



## (3) Unsupervised Data Methods

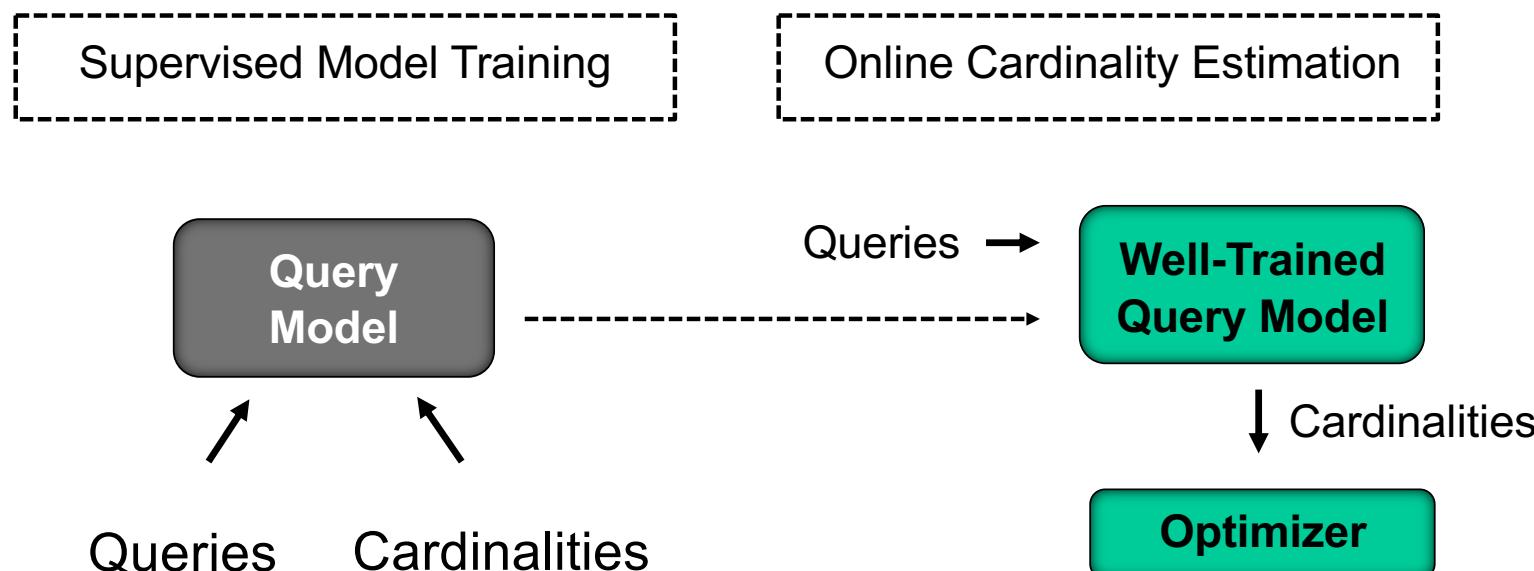
- Sum product network
- Autoregressive (AR) model
- Normalizing Flow (NF) model



# 1 Supervised Query Methods for Cardinality Estimation

## □ Problem Definition

A regression problem: learn the mapping function between query  $Q$  and its actual cardinality





# 1.1 Query-Driven: Neural Network on Single Tables

□ **Motivation:** Traditional estimation methods assume column independence.

□ **Solution:**

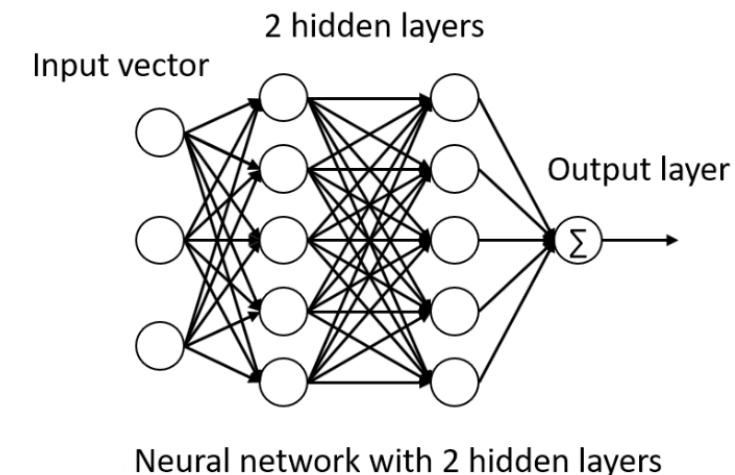
## ➤ Training

- Represent a query  $(lb_1 < A_1 < ub_1, \dots, lb_d < A_d < ub_d)$  on a table  $T$  with  $d$  attributes  $A_1, \dots, A_d$  as  $< lb_1, ub_1, \dots, lb_d, ub_d >$ .
- A neural network with two hidden layers is used to fit the mapping between the representation of the query and its cardinality.

## ➤ Inference

- Answer a query by the trained network.

$$(c_1 \leq lb_1 < c_2) \wedge (c_3 < ub_1 \leq c_4) \wedge (c_5 \leq ub_2 \leq c_6)$$





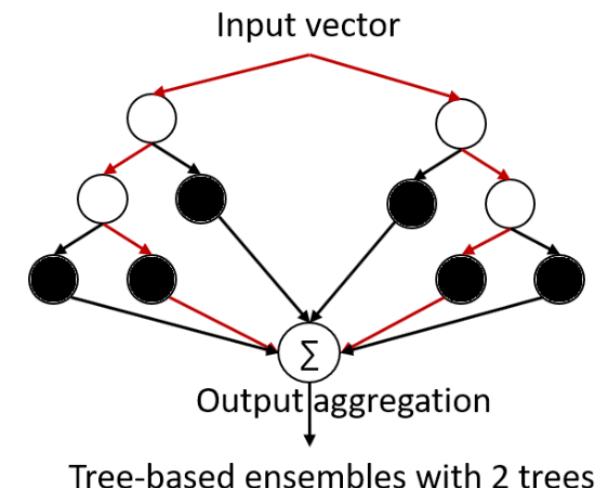
## 1.2 Query-Driven: XGBoost for Cardinality Estimation

### □ Solution:

- The query is **represented** in the same way as the neural network based approach.
- Use XGBoost, a decision tree-based ensemble model to **fit a mapping** between a query's representation and its cardinality.

### □ Comparison with Neural Network:

- Neural Network-based method are better **when training data is sufficient**.
- XGBoost is better when **the training data is insufficient**.



Tree-based ensembles with 2 trees

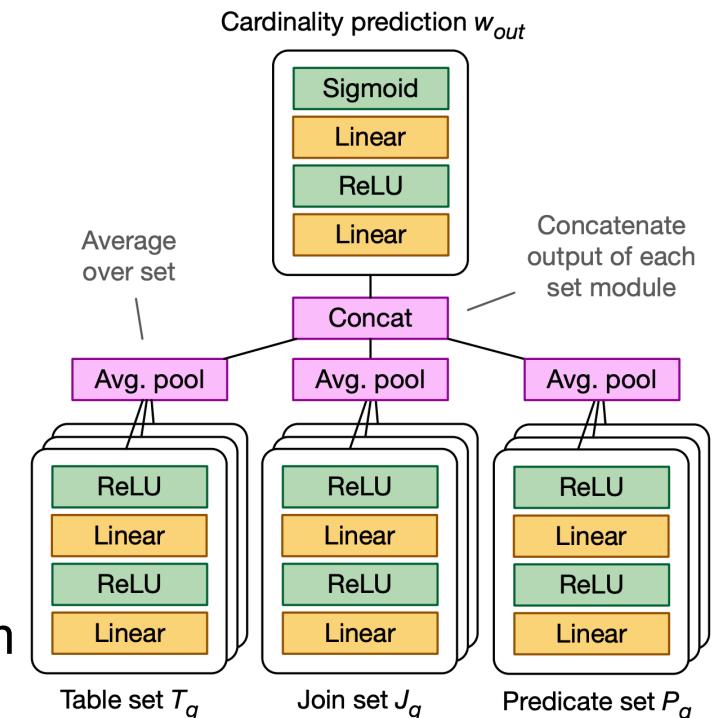


## 1.3 Query-Driven: Deep Learning for Multi-tables

□ **Motivation:** It's difficult for traditional methods to capture join-crossing correlations.

□ **Solution:**

- For **table set** and **join set** in the input, encode each table, join with one-hot encoding.
- For **predicates** of the form  $(\text{col}, \text{op}, \text{val})$ , encode columns **col** and operators **op** with one-hot encoding, and represent **val** as a normalized value in  $[0, 1]$ .
- Use some samples to address 0-tuple problem





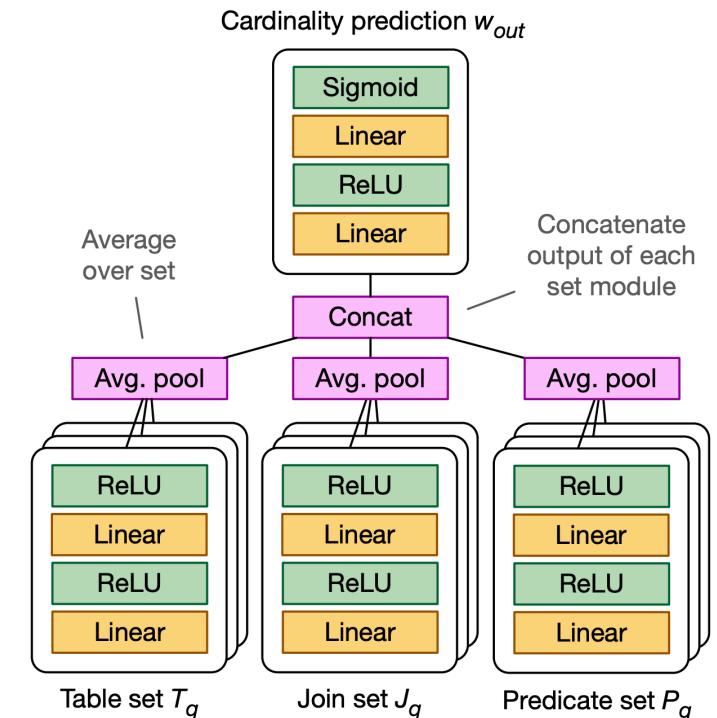
## 1.3 Query-Driven: Deep Learning for Multi-tables

### □ Training:

- The three parts of the input are spliced together after going through the linear layer, activation layer, and the dimensionality reduction layer.
- Then go through the linear and activation layer again to get the estimated cardinality.
- Tuning model parameters by backpropagating gradients.

### □ Inference

- Answer a query by the trained network.

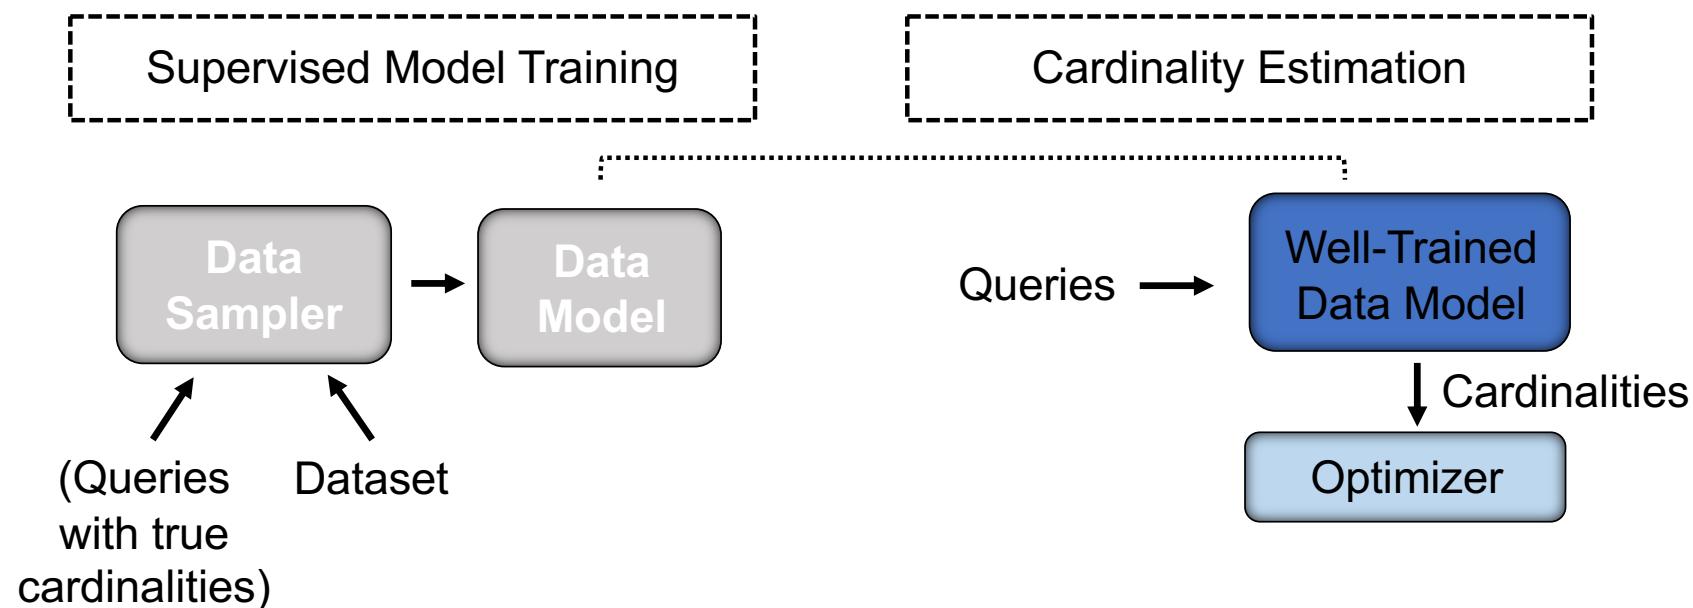




## 2 Supervised Data Methods for Cardinality Estimation

### □ Problem Definition

A **density estimation problem**: learn a joint data distribution of each data point. (except for the pre-training summarization model)





## 2.1 Supervised Data-Driven: Kernel-Density Model on Single Table

- **Motivation:** Multi-dimensional histograms are complex to construct and hard to maintain.
- **Key-idea:** Fit the probability density distribution of a data table by kernel density model.

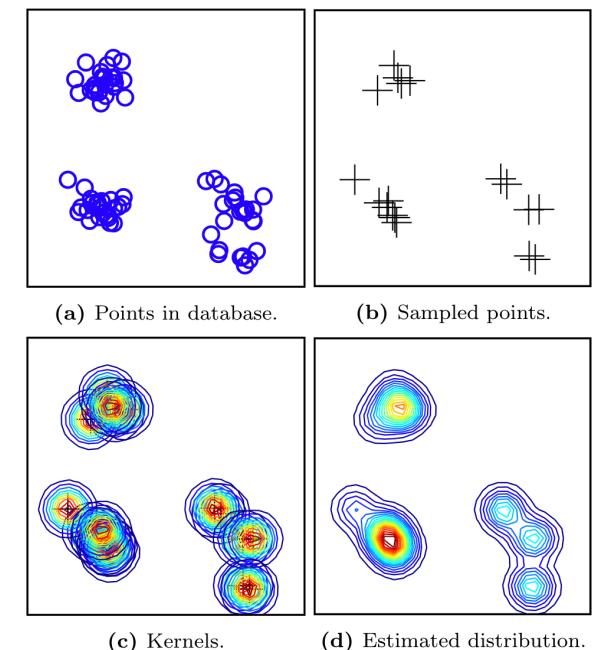
### ➤ Kernel-Density Model

**Model:**  $\hat{p}_H(\vec{x}) = \frac{1}{s \cdot |H|} \sum_{i=1}^s K\left(H^{-1}\left[\vec{t}^{(i)} - \vec{x}\right]\right)$

- $s$  is sample size;  $K$  is Gaussian function;
- $\vec{t}$  is sampled point;  $H$  is a parameter that needs to be learned.,

➤ **Inference:**  $\hat{p}_H(\Omega) = \int_{\Omega} \hat{p}(\vec{x}) d\vec{x}$

- $\Omega$  is the space represented by a query.





## 2.1 Supervised Data-Driven: Kernel-Density Model on Single Table

### Kernel-Density Estimation

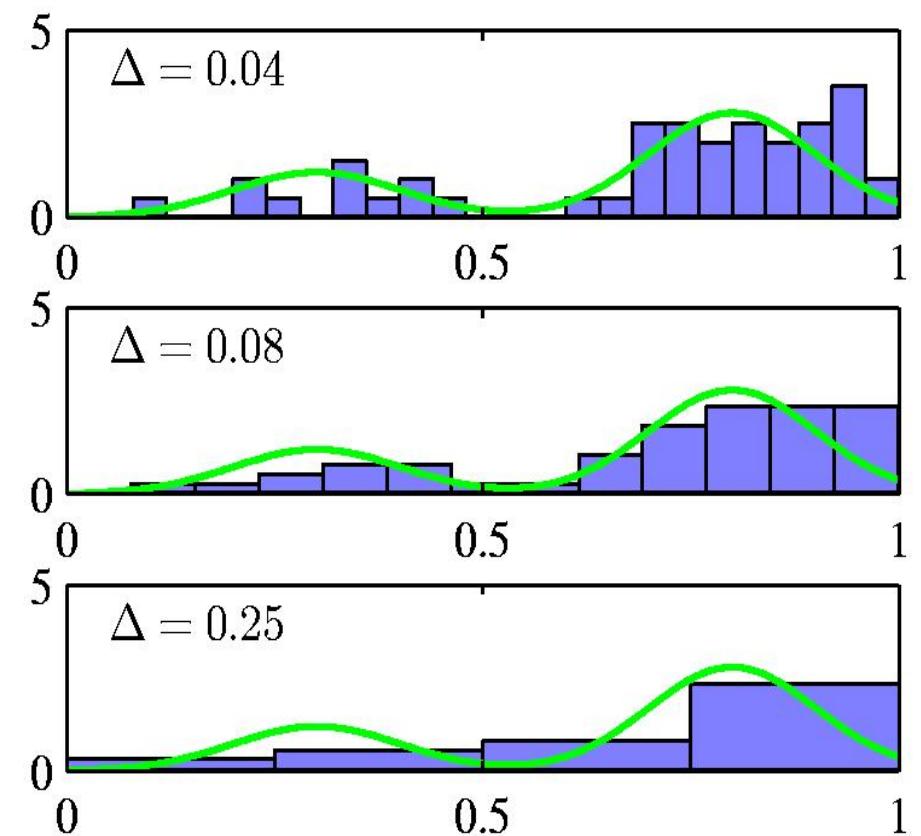
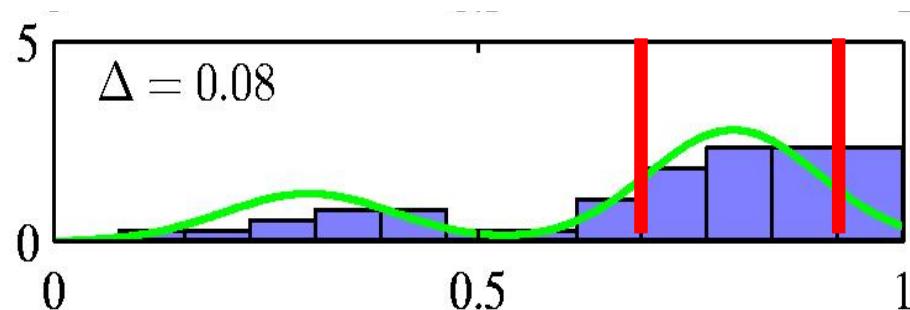
$$p_i = \frac{n_i}{N \Delta_i}$$

$p_i$ : probability

$n_i$ : number of points

$N$ : total number of points

$\Delta$ : bandwidth





## 2.1 Supervised Data-Driven: Kernel-Density Model on Single Table

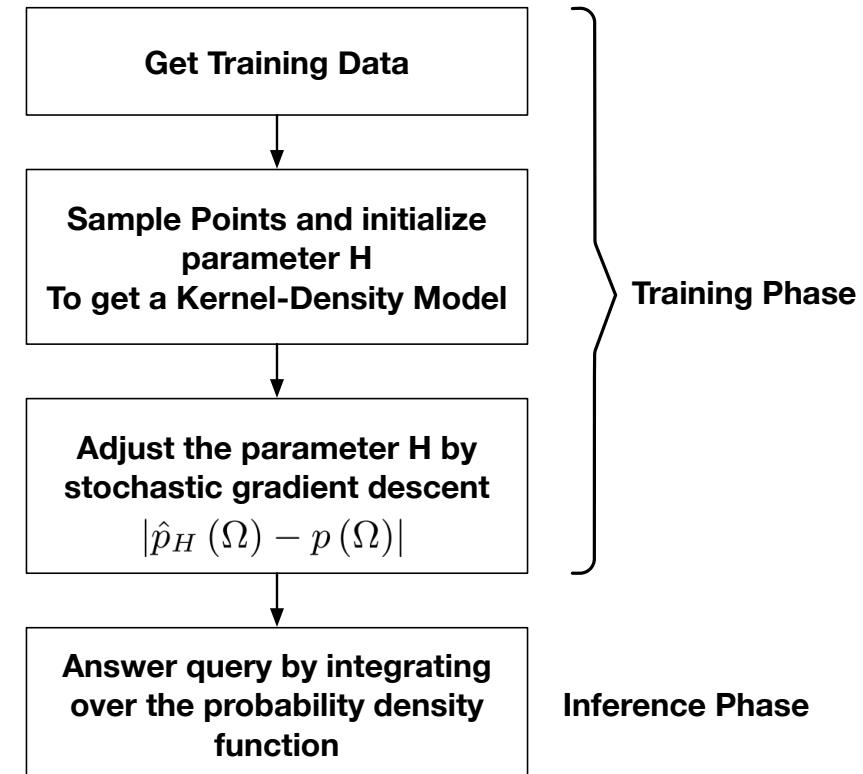
### □ Solution

#### □ Training

- Get a lot of queries with true cardinalities.
- Sample points (rows) from the table and initialize the parameter  $H$ .
- Adjust the parameter  $H$  by stochastic gradient descent according to estimated cardinalities by Kernel-Density Model and the true cardinalities.

#### □ Inference

- Answer queries by accumulating kernel density based on the kernel-density model.





## 2.2 Supervised Data-Driven: Uniform Mixture Model on Single Table

- **Motivation:** Traditional methods need to be populated in advance by performing costly table scans.
- **Key-idea:** Fit the probability density distribution of a data table by uniform mixture model.

### □ Uniform Mixture Model:

**Model:**  $f(x) = \sum_{z=1}^m h(z) g_z(x) = \sum_{z=1}^m w_z g_z(x)$

**Inference:**  $\int_{B_i} f(x) dx = \int_{B_i} \sum_{z=1}^m w_z g_z(x) dx$

- $w_z$  is the weight for a subpopulation  $z$
- $g_z(x) = 1/|G_z|$  is uniform distribution function
- $|G_z|$  is area of subpopulation  $z$

- $B_i$  is the space represented by a query.



## 2.2 Supervised Data-Driven: Uniform Mixture Model on Single Table

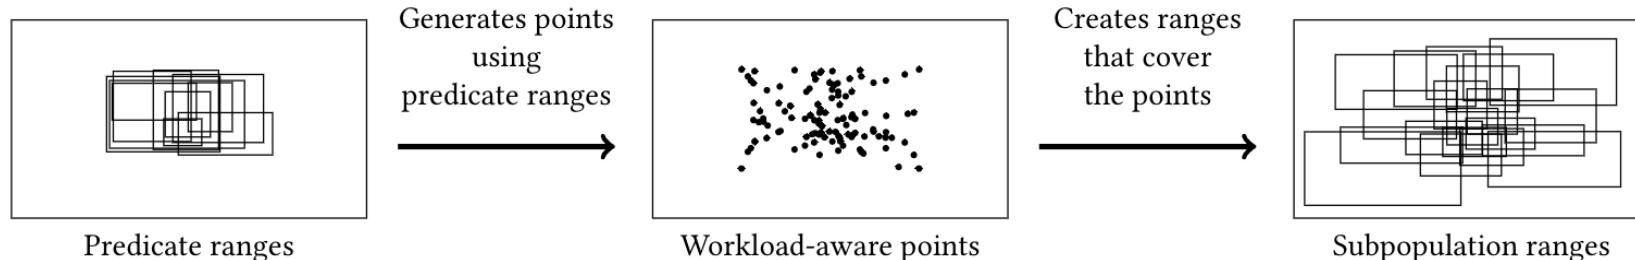
### □ Solution:

#### □ Training

- Sample some points within queries with true cardinalities.
- Generating subgroups for the points.
- Learn the weights  $w_z$  of the uniformity mixture model.

#### □ Inference

- Answer a query by Calculate the cumulative probability density (i.e., selectivity) according to the mixture density function.
- overlap area between query rectangle and data rectangle





## 2.3 Supervised Data-Driven: Pre-training Summarization Model on Single Table

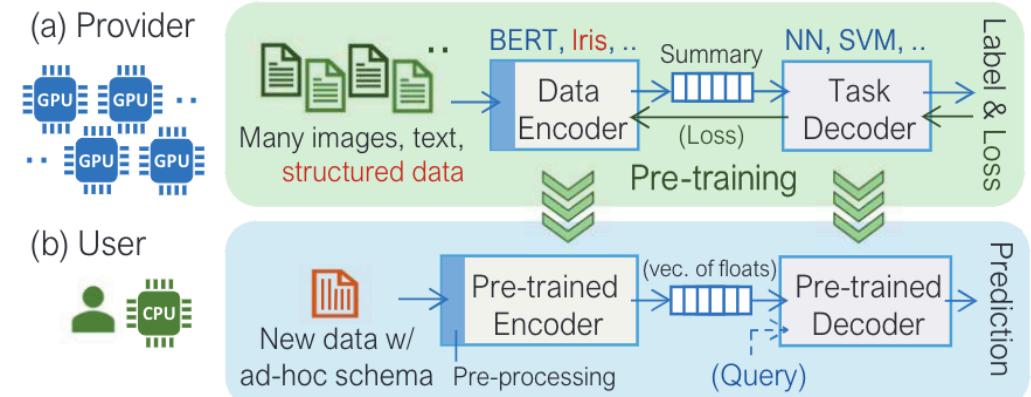
□ **Motivation:** Pre-training models avoid per-dataset training.

□ **Solution:**

□ Pre-train encoder and decoder with large data tables via gradient descent (Loss function is  $\left| \log \frac{\text{true card}}{\text{est. card}} \right|$ ).

□ Encode a table with the pre-trained encoder.

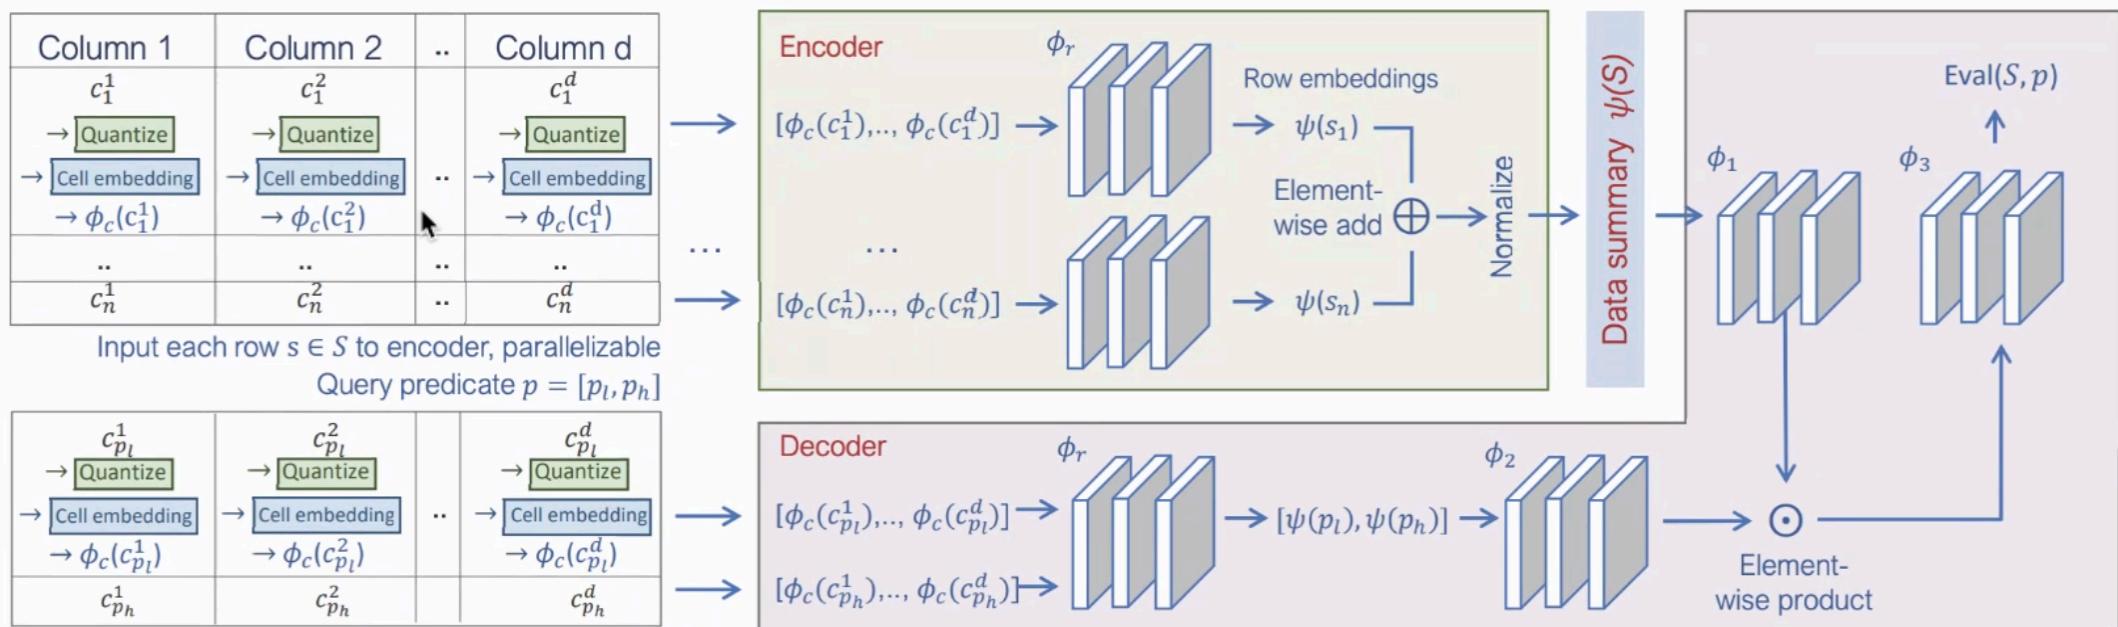
□ Input a query and the encoded result of the table into the pre-trained decoder to get the cardinality.





## 2.3 Supervised Data-Driven: Pre-training Summarization Model on Single Table

- Data Encoder
- Query Decoder

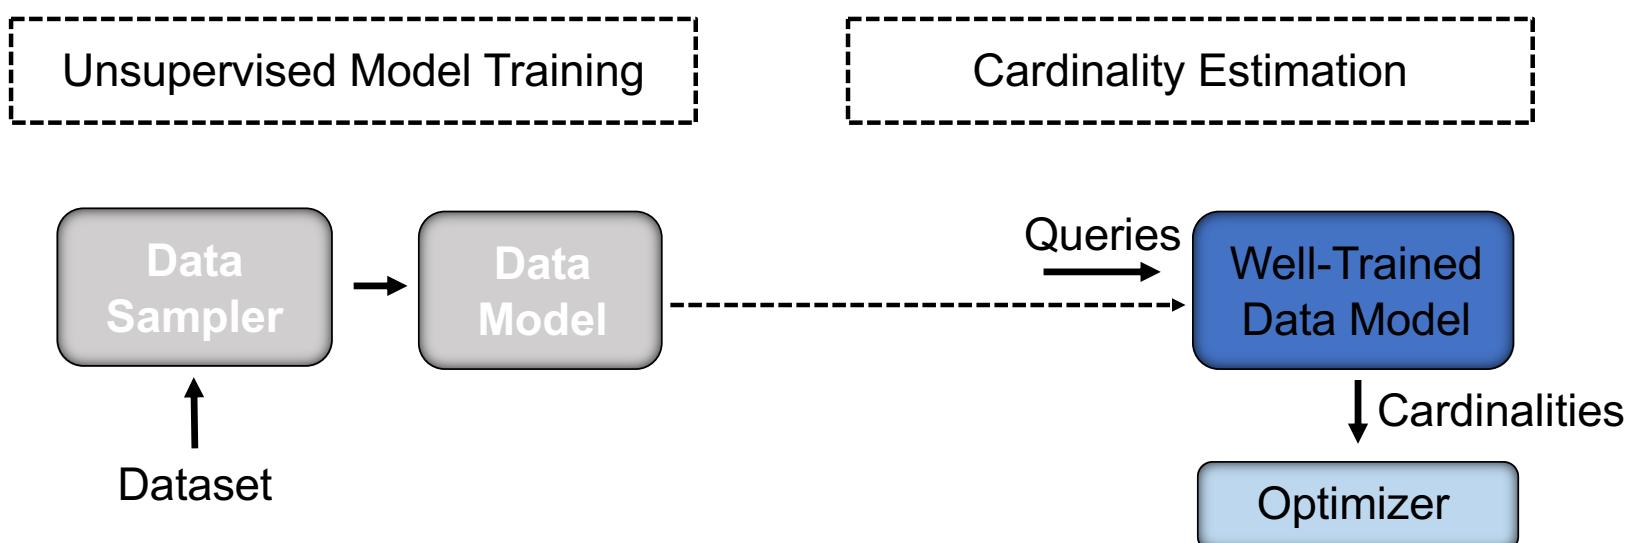




# 3 Unsupervised Data Methods for Cardinality Estimation

## □ Problem Definition

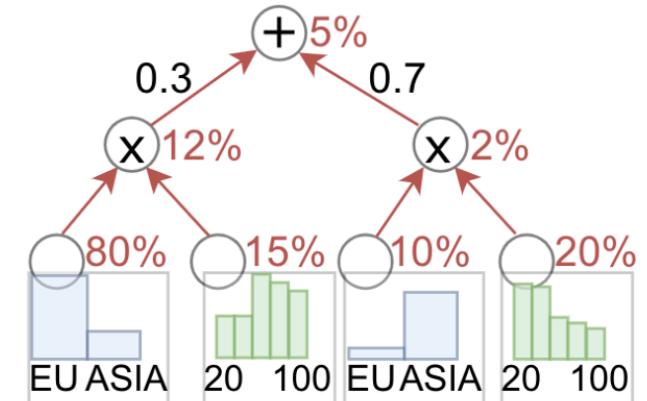
A **regression problem**: learn a probability function for each data point.





## 3.1 Unsupervised Data-Driven: Sum-Product Network for Multi-tables

- **Motivation:** Most of the existing estimators require SQL queries.
- **Base-idea:** Learn the joint probability distribution by Sum-Product Network .
- **Relational Sum Product Network (RSPN)**
  - RSPN consists of three types of node:
    - product node: split the columns of a table.
    - sum node: split the rows of a table.
    - leaf node: represent probability distributions for individual variables.



(d) Probability of European Customers younger than 30



## 3.1 Unsupervised Data-Driven: Sum-Product Network for Multi-tables

### □ Solution:

#### □ Training

- Generate some queries and their cardinalities.
- Build a RSPN by recursively partitioning
  - Row: K-means clustering
  - Column: randomized dependency coefficient.

#### □ Inference

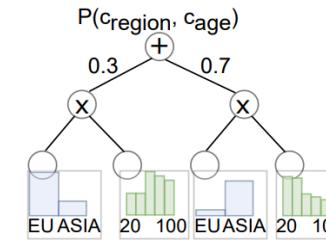
- Estimate cardinality in **bottom-up**.
- Supports multi-table queries via join sampling, and supports queries on sub-schemas based on **fanout scaling**.

c_id	c_age	c_region
1	80	EU
2	70	EU
3	60	ASIA
4	20	EU
...	...	...
998	20	ASIA
998	25	EU
999	30	ASIA
1000	70	ASIA

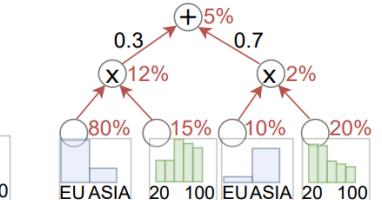
(a) Example Table

c_age	c_region
80	EU
70	EU
60	ASIA
20	EU
...	...
...	...
20	ASIA
25	EU
30	ASIA
70	ASIA

(b) Learning with Row/Column Clustering



(c) Resulting SPN



(d) Probability of European Customers younger than 30



# Fanout scaling for multi-tables

**Fanout scaling** is to support sub-schema queries on a full outer join table. (full outer join table contains duplicate tuples)

- **Solution:** for each *foreign key* → *primary key* relationship add a column denoting how many corresponding join partners a tuple has.

- **Example:**

```
SELECT COUNT(*)
  FROM CUSTOMER C
 WHERE c_region='EUROPE' ;
```

- True answer is 2, but there are 3 tuples in the outer join table.

- Fanout scaling result:  $|C \bowtie O| \cdot \mathbb{E}(1/\mathcal{F}'_{C \leftarrow O} \cdot \mathbf{1}_{c\_region='EU'}} \cdot \mathcal{N}_C)$

$$= 5 \cdot \frac{1/2+1/2+1}{5} = 2$$

Customer			Order		
c_id	c_age	c_region	o_id	c_id	o_channel
1	20	EUROPE	1	1	ONLINE
2	50	EUROPE	2	1	STORE
3	80	ASIA	3	3	ONLINE
			4	3	STORE

Customer			Order				
$\mathcal{N}_C$	c_id	c_age	c_region	$\mathcal{F}'_{C \leftarrow O}$	$\mathcal{N}_O$	o_id	o_channel
1	1	20	EUROPE	2	1	1	ONLINE
1	1	20	EUROPE	2	1	2	STORE
1	2	50	EUROPE	1	0	NULL	NULL
1	3	80	ASIA	2	1	3	ONLINE
1	3	80	ASIA	2	1	4	STORE

## 3.2 Unsupervised Data-Driven: Autoregressive Model on Single Table

□ **Motivation:** Existing estimators struggle to capture the rich multivariate distributions of relational tables.

□ **Solution:**

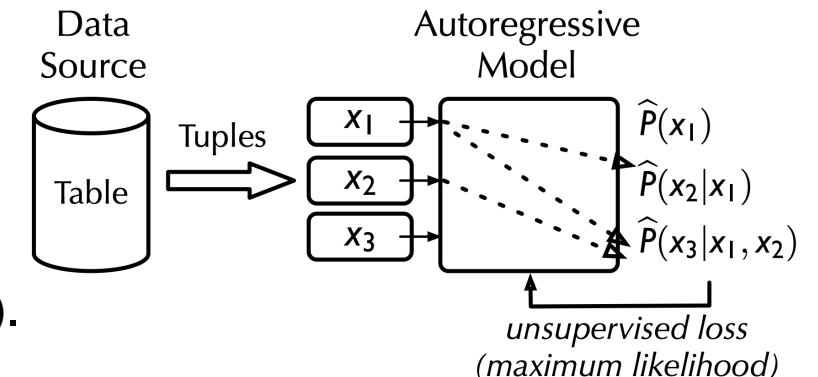
□ **Training:** Use Autoregressive (AR) Model to fit the joint probability of different columns.

□ **Inference:** Estimate the cardinality of the equivalent query based on the AR model.

□ Monte Carlo sampling

□ Range queries are supported by progressive sampling (sampling by learned probability distribution).

$$\begin{aligned}\hat{P}(\mathbf{x}) &= \hat{P}(x_1, x_2, \dots, x_n) \\ &= \hat{P}(x_1)\hat{P}(x_2|x_1)\cdots\hat{P}(x_n|x_1, \dots, x_{n-1})\end{aligned}$$



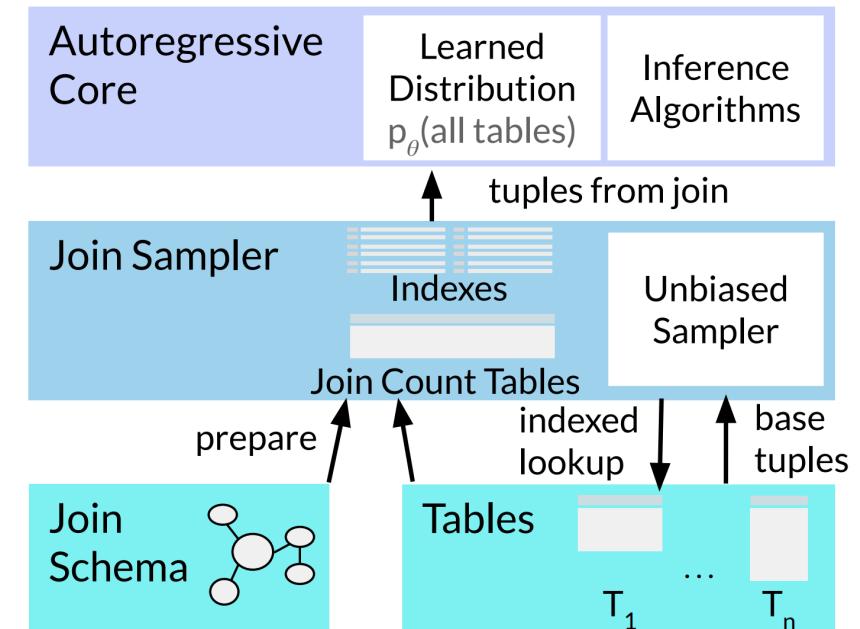


## 3.2 Unsupervised Data-Driven: Autoregressive Model on Single Table

- **Motivation:** Previous AR models do not support multi-table queries.

- **Solution:**

- Learn an autoregressive model for the **outer join of all tables**.
- Supports multi-table queries via join sampling, and supports queries on sub-schemas through **fanout scaling**.



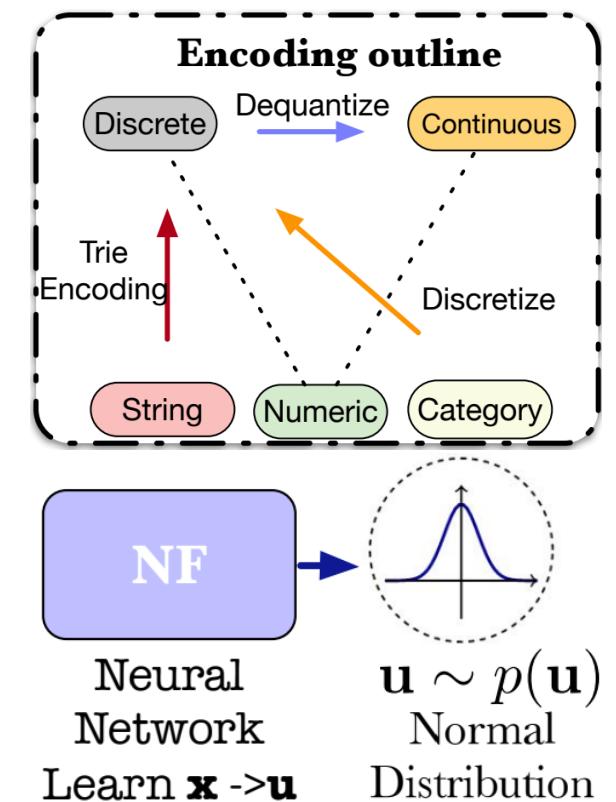


### 3.3 Unsupervised Data-Driven: Normalizing Flow (NF) model for Multi-tables

□ **Motivation:** Previous data-driven approaches do not handle tables with large **domain sizes** well.

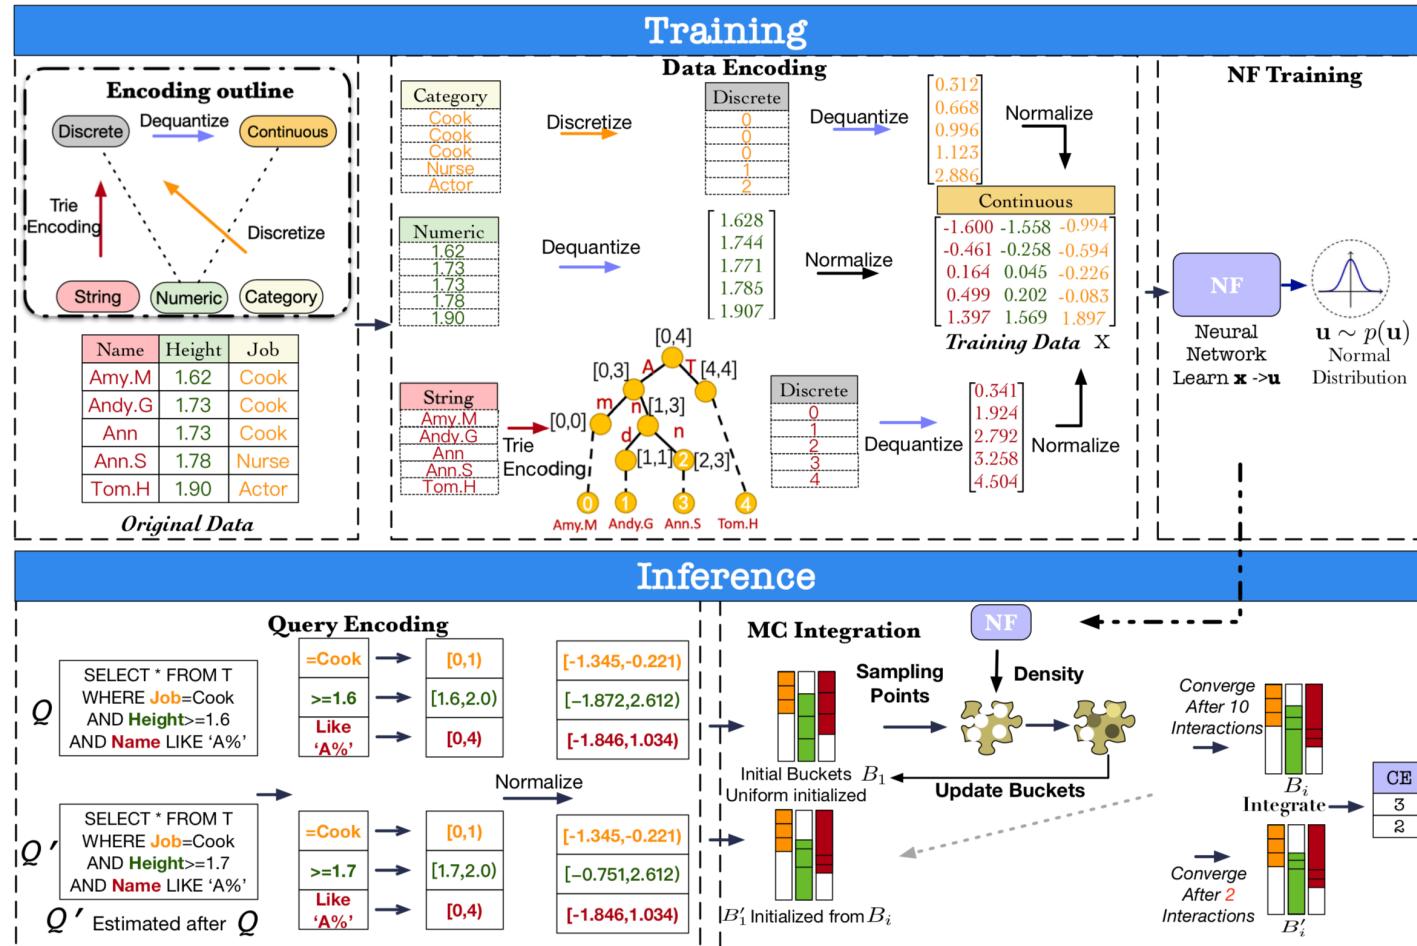
□ **Solution:**

- Dequantize and **normalize** discrete variables to **continuous** variables.
- Use **Normalizing Flow (NF)** model to learn the joint probability distribution of data points.
- Accumulating continuous normalized flow distribution function by adaptive importance sampling to answer a query.
- Support multi-table query through **fanout scaling**.





### 3.3 Unsupervised Data-Driven: Normalizing Flow (NF) model for Multi-tables





# Summarization of Learned Cardinality Estimation

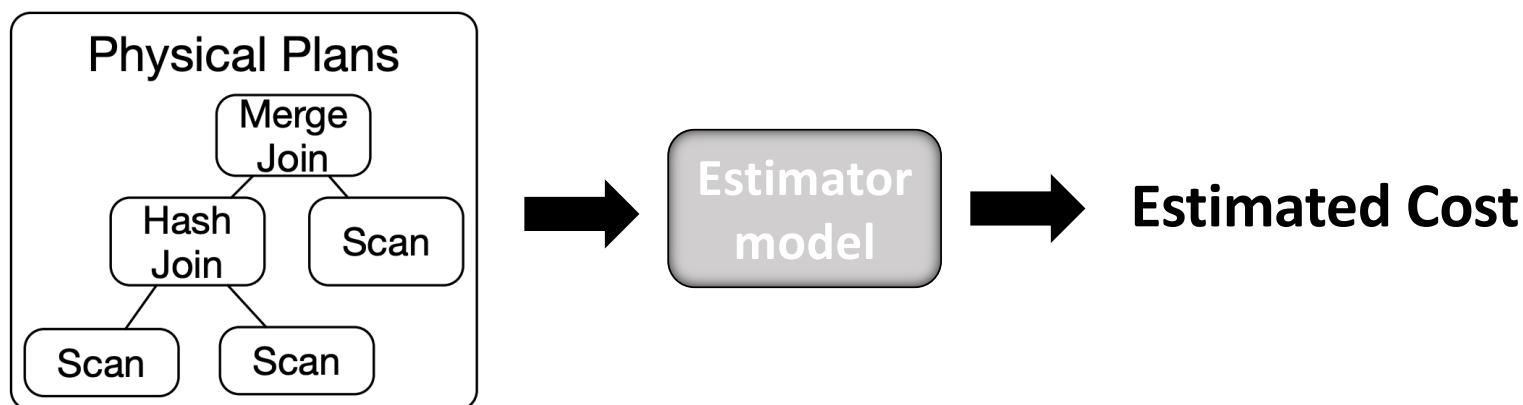
Method	Quality	Training overhead	Training Data	Adaptive	Model Size	Inference Latency
Lightweight Neural Network & XGBoost	✓	low	many queries	✓	small	fast
Convolutional Neural Network	✓✓	low	many queries	✓✓	small	fast
Kernel-Density Model	✓	medium	Data samples	✓	small	medium
Uniform Mixture Model	✓	medium	Data samples	✓	small	medium
Autoregressive (AR) Model	✓✓	high	Data samples	✓✓✓	high	slow
Sum-Product Network	✓✓	high	Data samples	✓✓✓	medium	medium
Normalizing Flow (NF) model	✓✓	high	Data samples	✓✓✓	medium	medium
Pre-training summarization model	✓	high	Lots of tables	✓✓	very small	fast



# Cost Estimation

## □ Problem Formulation:

- **Cost:** Execution cost of a query plan.
- **Input:** A SQL Query Plan.
- **Output:** An Estimated Cost.





# Relations of Cardinality/Cost Estimation

## ○ Task Target

- Cost estimation is to approximate the execution-time/resource-consumption.

## ○ Correlations

- Cost estimation is based on cardinality.

## ○ Estimation Difficulty

- Cost is harder to estimate than cardinality, which considers multiple factors (e.g., seq scan cost, CPU usage).

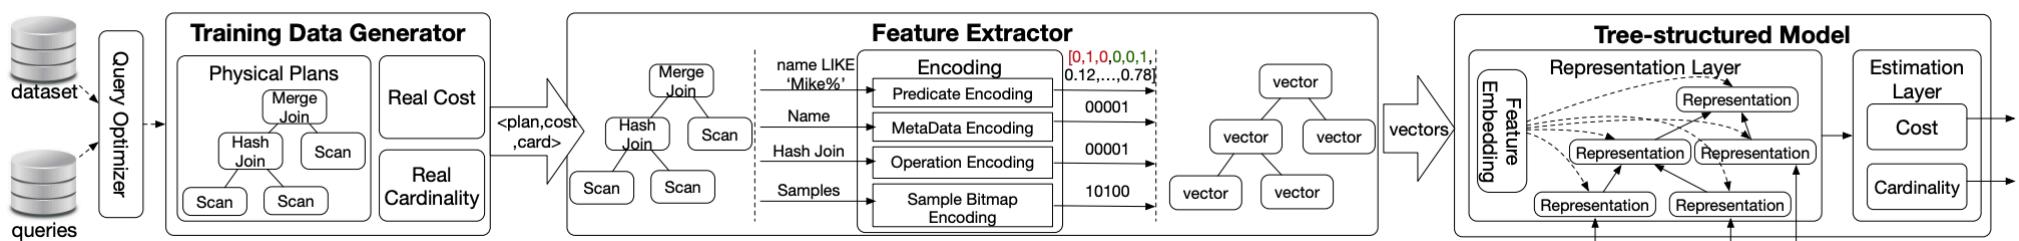


# Tree-LSTM for Cost Estimation

□ **Motivation:** Traditional cost estimation is inaccurate without learned plan representation.

□ **Solution:**

- Generate many query plans and true costs as training data.
- Encode the query plan via one-hot encoding.
- Representation layer learns an **embedding** of each query plan by Tree-LSTM.
- Estimation layer outputs estimated cost based on the representation layer's output.

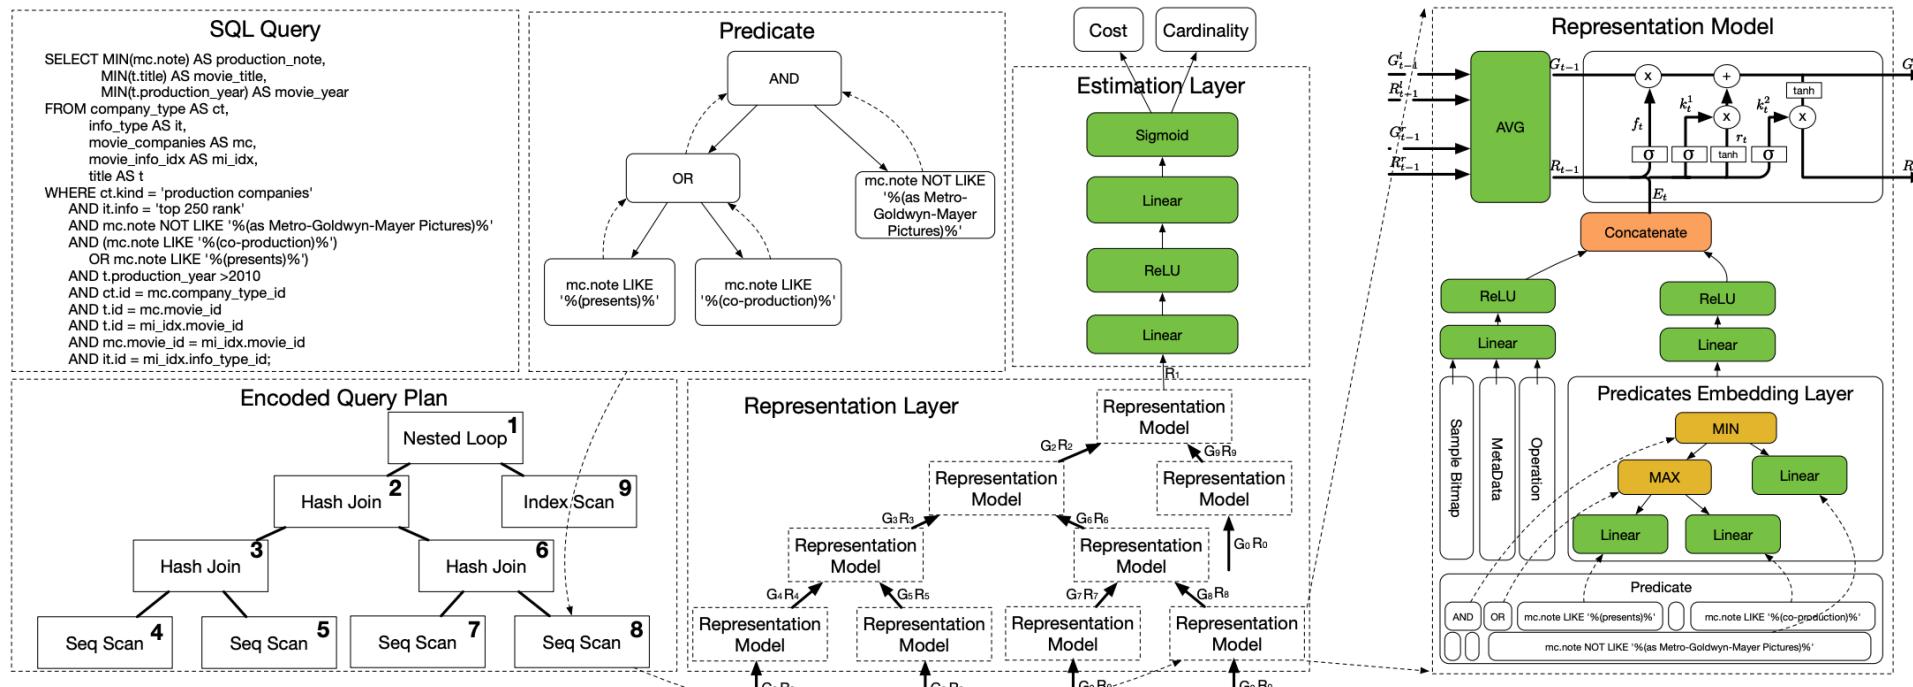




# Tree-LSTM for Cost Estimation

## □ Model Construction

- Traditional cost estimation uses estimated card, which is inaccurate without predicate encoding →

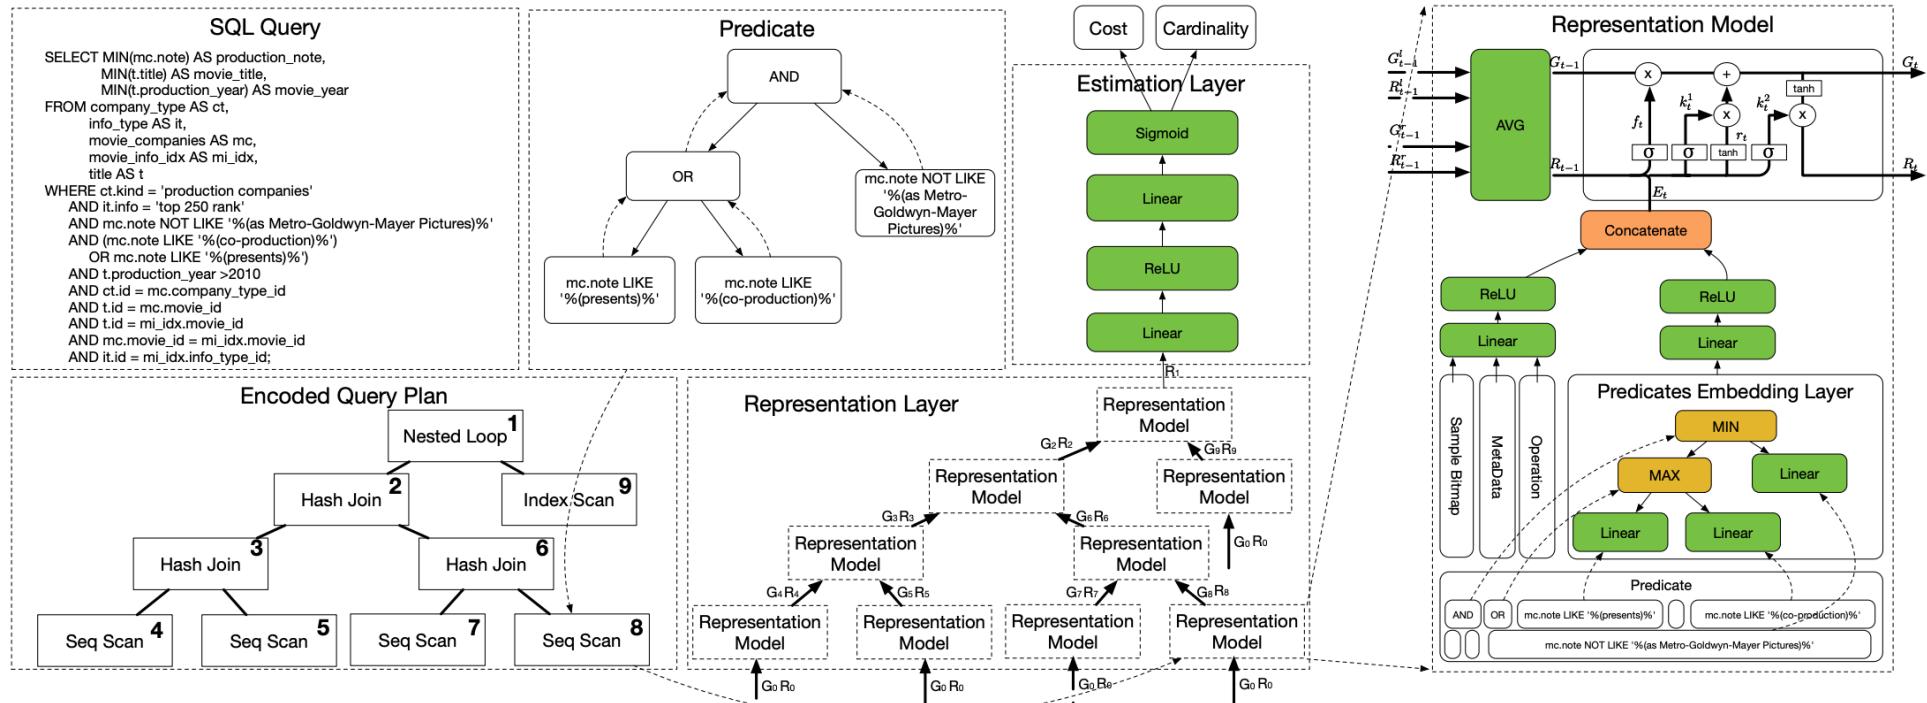


J. Sun and G. Li. An end-to-end learning-based cost estimator. PVLDB, 13(3):307–319, 2019.



# Tree-LSTM for Cost Estimation

- The representation layer learns an embedding of each subquery (global vector denotes the subquery, local vector denotes the root operator)
- The estimation layer outputs cardinality & cost simultaneously





# Take-aways of Cardinality Estimation

- Cardinality Estimation
  - Data-driven methods are more effective for single tables.
  - Query-driven methods are more efficient than Data-driven methods.
  - Data-driven methods are more robust than Query-driven methods.
  - Samples are crucial to most Data-driven methods.
- Cost Estimation
  - Accurate cost estimation requires better plan embedding.
- Open Problems
  - High Accuracy with small model size and inference latency
  - Adaptivity



# Deep Learning for Query Latency Estimation

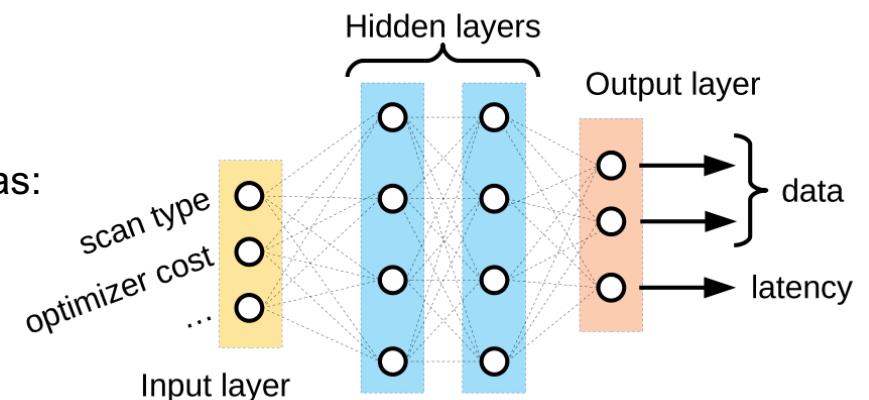
## □ Motivation

- Statistical methods fail to estimate based on query structures, and cause great errors;
- Compared with cost estimation, latency estimation is more complex because (1) it relies on system resources and (2) Cost is one important factor of latency estimation.

## □ Core Idea: Utilize deep learning to capture the relations between input tables, operators, and the final performance

## □ Solution

- Represent each operator  $q_i$  with a neural unit
  - E.g., for a scan operator, the neural unit is designed as:





# Deep Learning for Query Latency Estimation

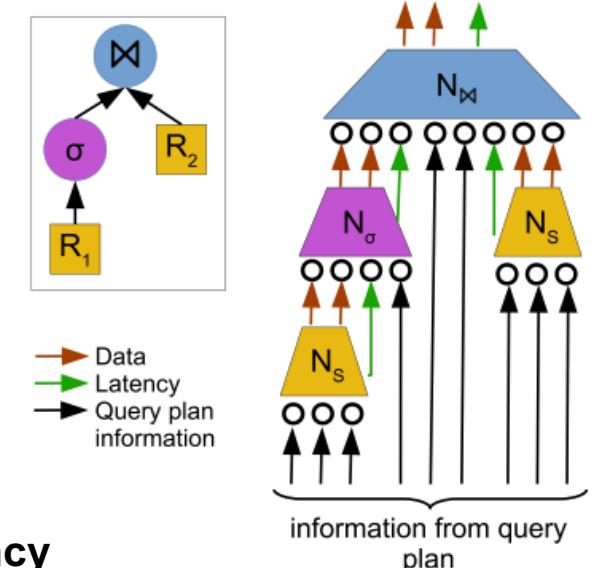
## □ Motivation

- Statistical methods fail to estimate based on query structures, and cause great errors;
- Compared with cost estimation, latency estimation is more complex because (1) it relies on system resources and (2) Cost is one important factor of latency estimation.

□ **Core Idea:** Utilize deep learning to capture the relations between input tables, operators, and the final performance

## □ Solution

- Represent each operator  $q_i$  with a neural unit
- Concatenate neural units by following the query structures
  - Example Query Q (2 Scans, 1 Join)
  - Tree-structured Network for Q:
    - The outputs of the scan units ( $N_\sigma + N_s$ )
    - Input of the join operator ( $N_{\bowtie}$ )
    - The final predicted query latency.
  - Matches the query structure to predict the query latency





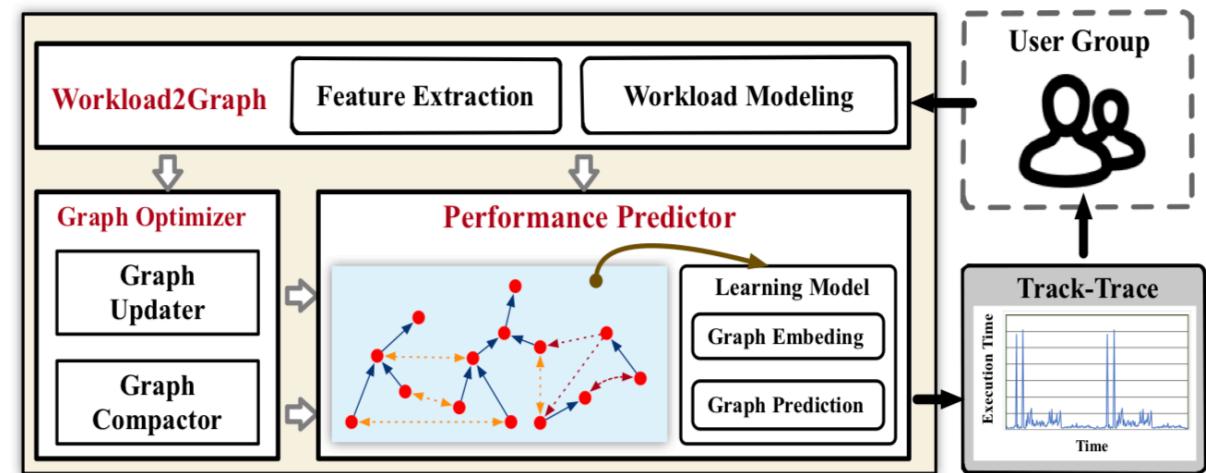
# Graph Embedding for Query Latency Estimation

## □ Latency Estimation for Concurrent Queries

- data-sharing
- data-conflict
- resource-competition
- parent-child relationship

## □ Graph-based method

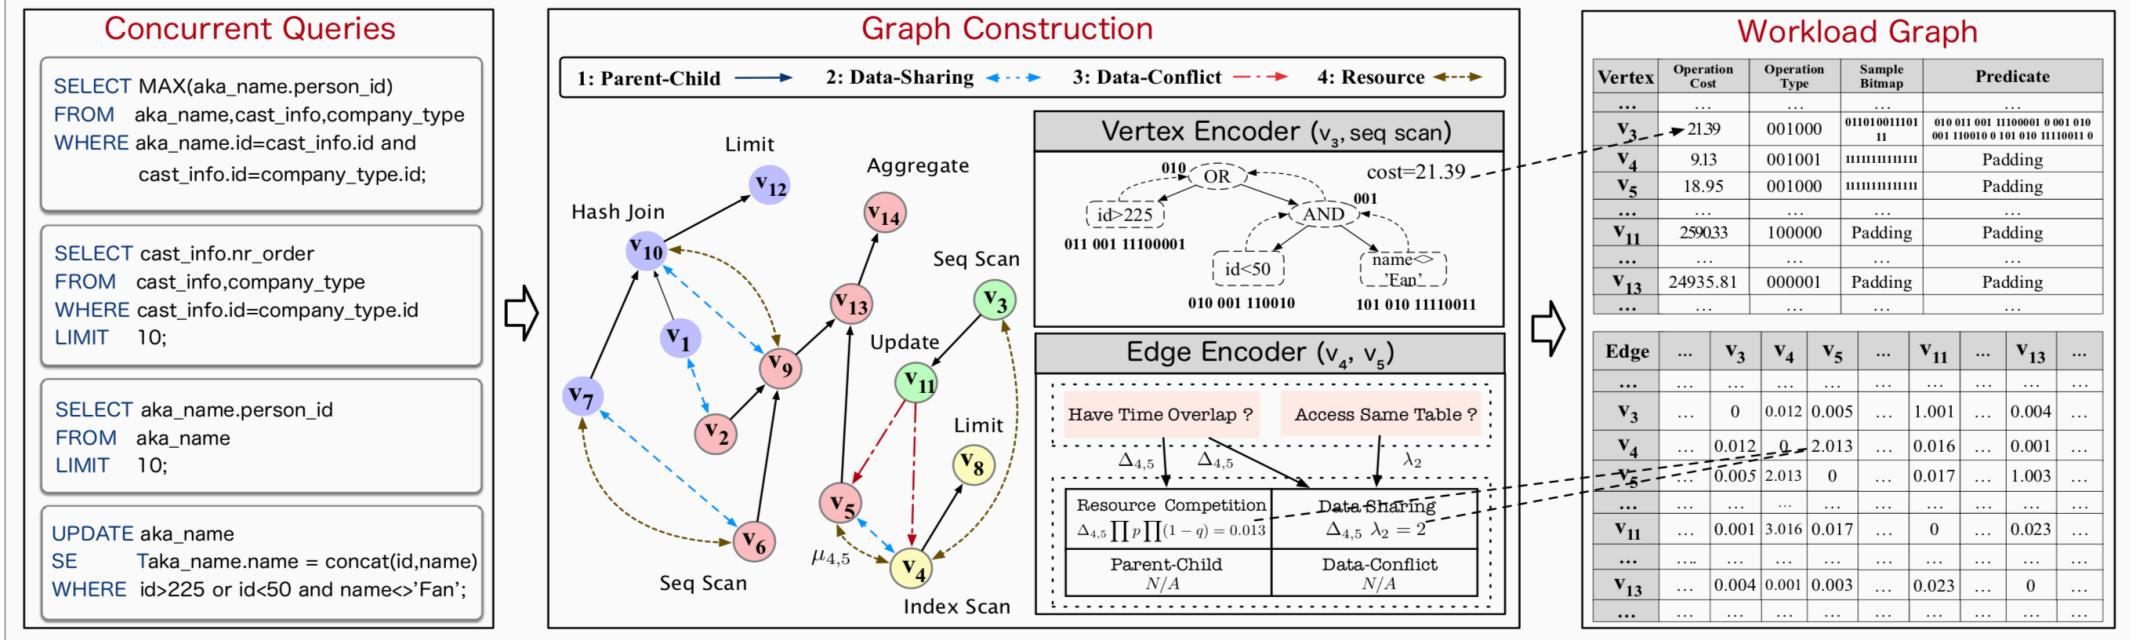
- **Workload2Graph:** graph modeling
- **Graph prediction:** GNN to predict the latency
- **Graph update:** on-the-fly update the model





# Graph Embedding for Query Latency Estimation

## Graph Modeling



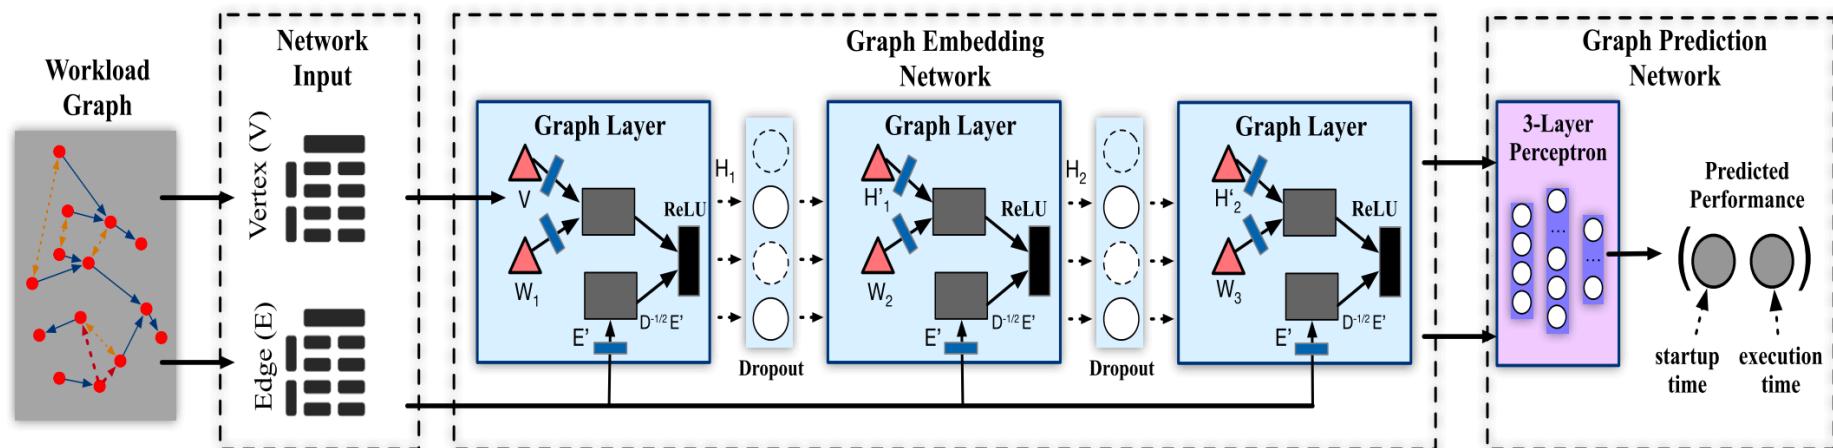


# Graph Embedding for Query Latency Estimation

## □ Model Construction

- **Performance prediction of concurrent queries**

- Represent concurrent queries with a graph model
- Embed the graph with graph convolution network and predict the latency of all the operators with a simple dense network





# Deep Learning for Index Benefit Estimation

## □ Challenge

### □ The index/view benefit is hard to evaluate

- Multiple evaluation metrics (e.g., index benefit, space cost)
- Cost estimation by the optimizer is inaccurate

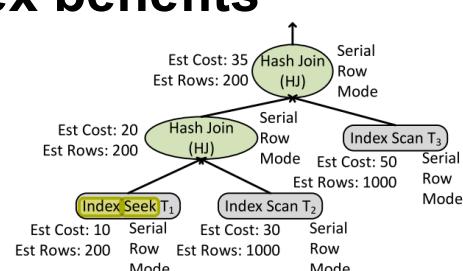
### □ Interactions between existing data structures

- Multiple column access, Data refresh
- Conflicts between MVs



# Deep Learning for Index Benefit Estimation

- Motivation: Critical to estimate index benefits by comparing execution costs of plans with/without created indexes
- Core Idea: Take benefit estimation as an ML classification task.
- Challenge: It is hard to accurately estimate the index benefits
- Solution:
  - Prepare training data
    - Query Plans + Costs under different indexes
  - Train the classification model
    - Input: Two query plans with/without indexes
    - Output: 1 denotes performance gains; 0 denotes no gains
  - Solve the index selection problem
    - Use the model to create indexes with performance gains



(a) Example query plan.

EstNodeCost	LeafWeightEstRows
Seek_Row_Serial	200
Scan_Row_Serial	2000
HJ_Row_Serial	4600
NLJ_Row_Serial	0
MJ_Row_Serial	0
...	...

EstNodeCost

LeafWeightEstRows  
WeightedSum

Seek_Row_Serial	10
Scan_Row_Serial	80
HJ_Row_Serial	55
NLJ_Row_Serial	0
MJ_Row_Serial	0
...	...

...

Seek_Row_Serial	200
Scan_Row_Serial	2000
HJ_Row_Serial	4600
NLJ_Row_Serial	0
MJ_Row_Serial	0
...	...

(b) Feature channels for the plan.



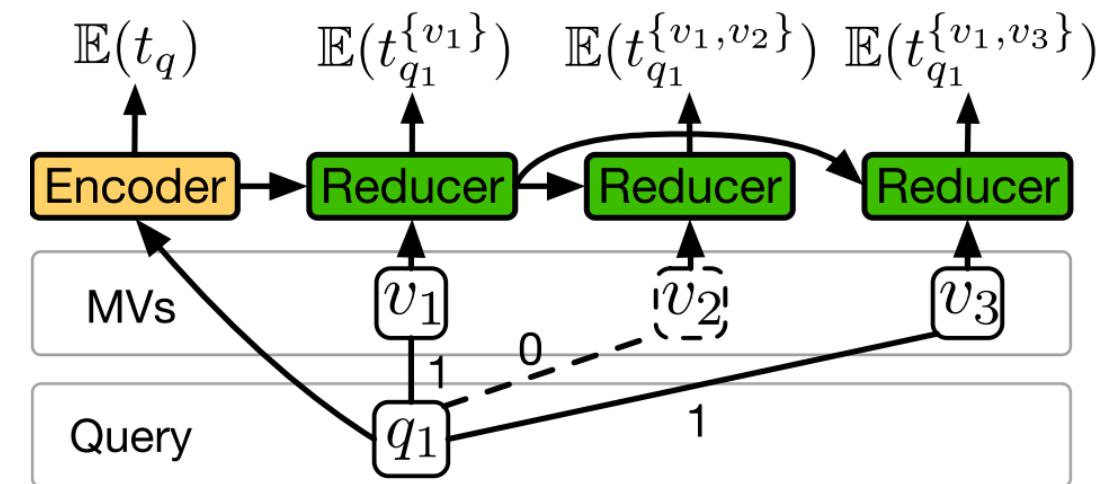
# Encoder-Decoder for View Benefit Estimation

## □ Feature Extraction

- Previous work take candidate views as fixed length →
- Encode various number and length of queries and views with an *encoder-reducer model*, which captures correlations with attention

## □ Model Construction

- It is hard to jointly consider MVs that may have conflicts →
- (1) Split the problem into sub-steps that select one MV; (2) Use attention-based model to estimate the MV benefit



Y. Han, G. Li, H. Yuan, and J. Sun. An autonomous materialized view management system with deep reinforcement learning. In ICDE, 2021.



## Take-aways of Benefit Estimation

- Learned utility estimation is more accurate than traditional empirical methods
- Learned utility estimation is also accurate for multiple-MV optimization
- Query encoding models need to be trained periodically when data update
- Open problems:
  - Benefit prediction for future workload
  - Cost of initialization and future updates

Thanks

